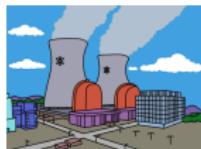
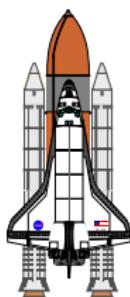


# SuperGlue: IDL-Based, System-Level Fault Tolerance for Embedded Systems

June 18, 2016

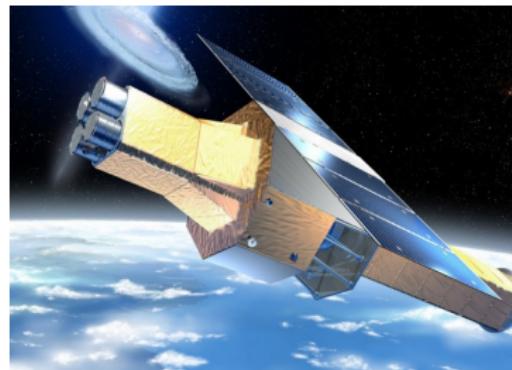
Computer Science Department  
The George Washington University

# Real-Time and Embedded Systems

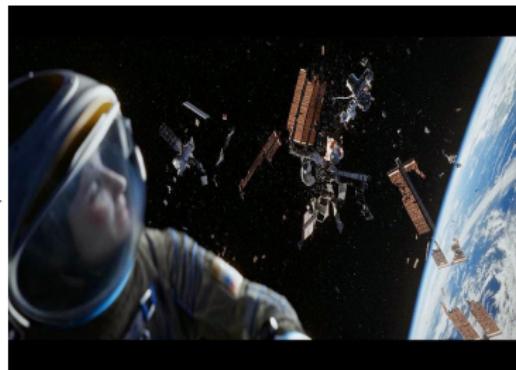


# Consequences of Embedded System Faults

Japanese X-ray astronomy satellite Hitomi lost in 2016



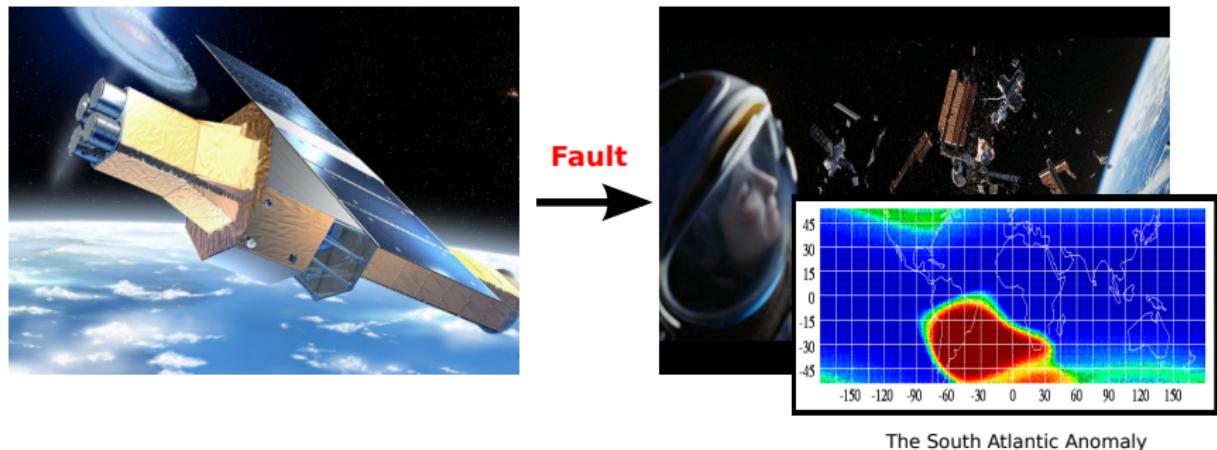
Fault  
→



- Financial losses of \$286 million USD
- “*It's a scientific tragedy*” - Richard Mushotzky, UMD

# Consequences of Embedded System Faults

Japanese X-ray astronomy satellite Hitomi lost in 2016



- Financial losses of \$286 million USD
- “*It’s a scientific tragedy*” - Richard Mushotzky, UMD

# Consequences of embedded system faults

## Toyota Sudden Unintended Acceleration (SUA) in 2004 – 2010



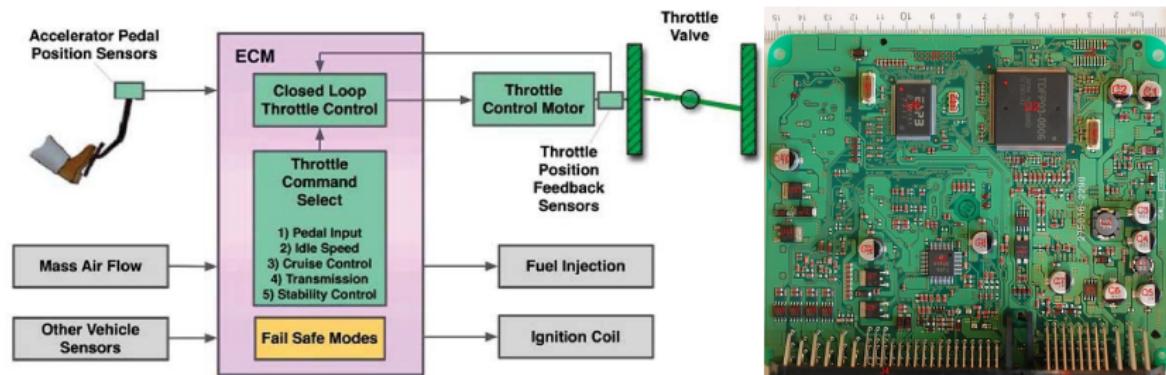
Fault →



- 89 deaths as of May 2010 and nearly 400 U.S. lawsuits
- Recall 10+ million vehicles and pay \$1.2 Billion USD fine

<http://www.cbsnews.com/news/toyota-unintended-acceleration-has-killed-89/>

# Sudden Unintended Acceleration

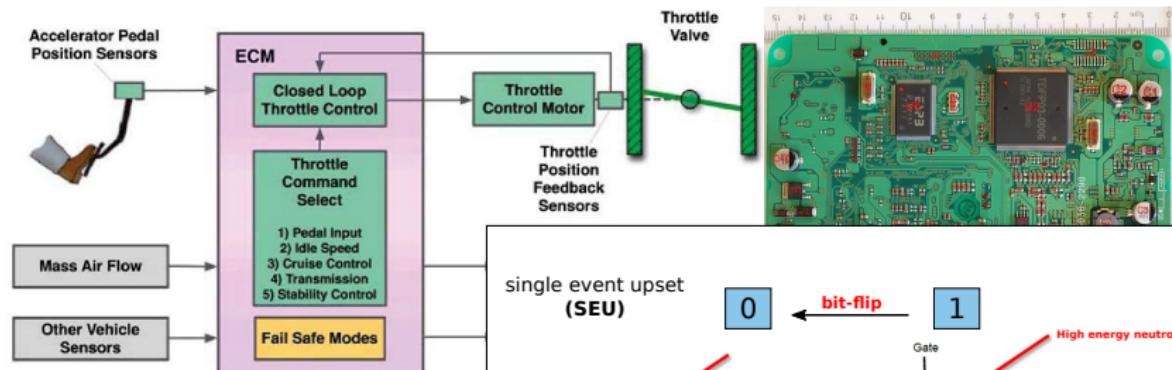


## Uncontrolled acceleration in Toyota Camrys

- Electronic Throttle Control System (ETCS)
- OSEK OS, 24 tasks, 280K LOC of C
- bit flip in scheduler data-structures  
→ reproducible 30-sec uncontrolled acceleration

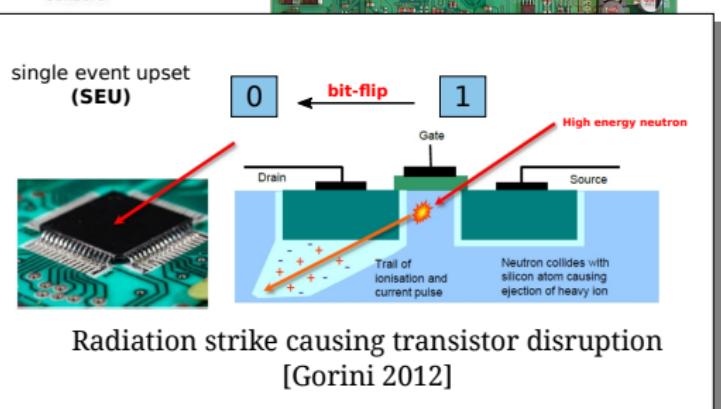
*“a bit-flip there, will have the effect of killing one of the tasks”*  
– Bookout v Toyota

# Sudden Unintended Acceleration



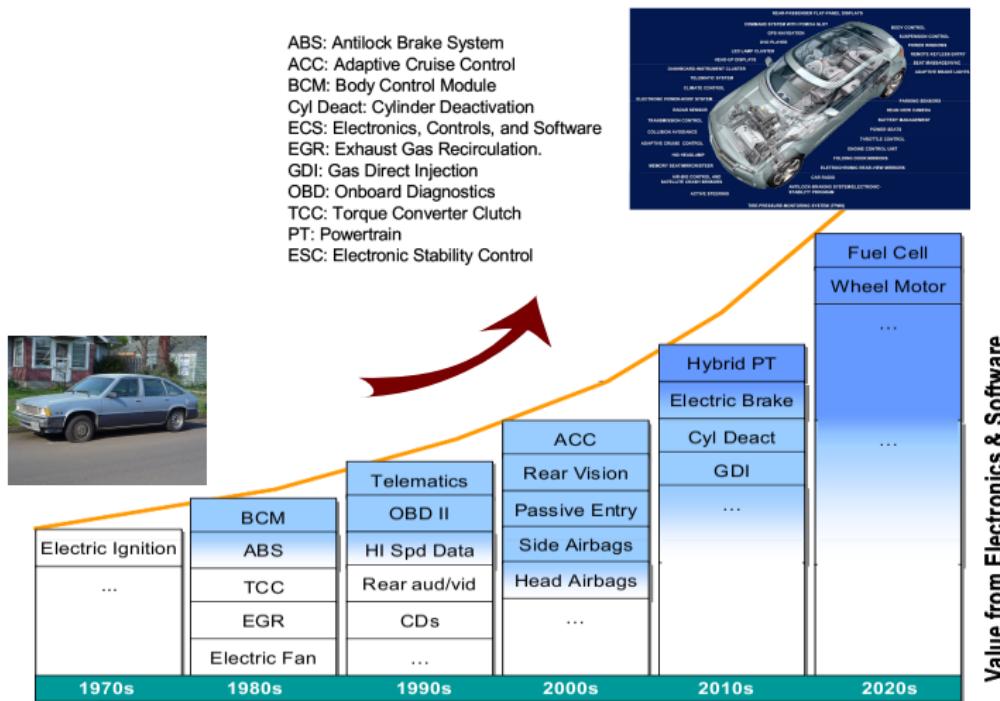
Uncontrolled acceleration in

- Electronic Throttle Con
- OSEK OS, 24 tasks, 28
- bit flip in scheduler data-structures  
→ reproducible 30-sec uncontrolled acceleration



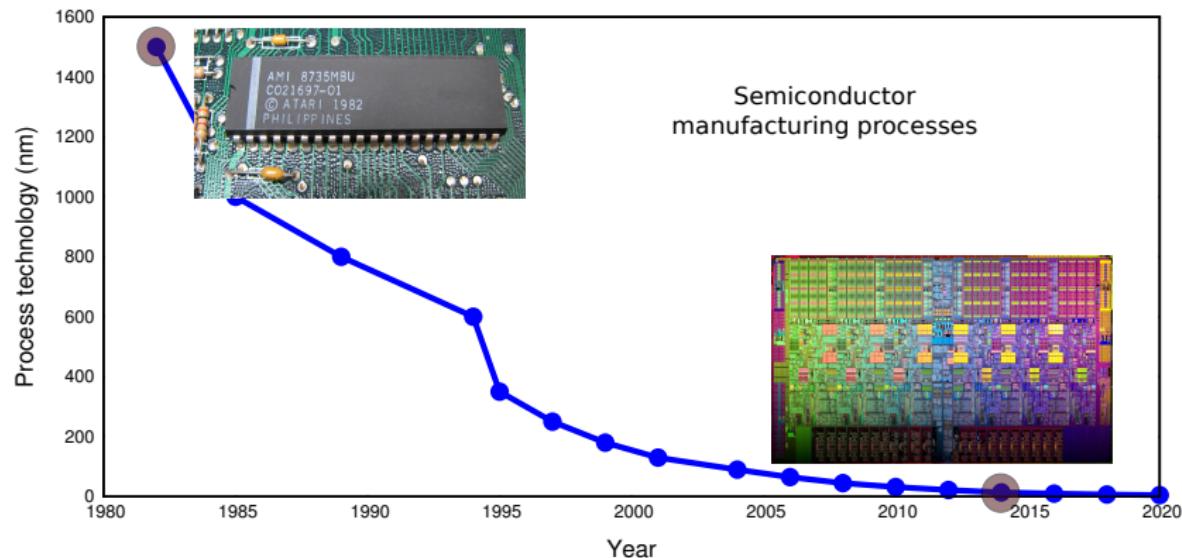
*"a bit-flip there, will have the effect of killing one of the tasks"*  
– Bookout v Toyota

# Embedded faults: Bad Now, Worse Tomorrow



- + more functionality
- more complexity → dependability more challenging

# Embedded faults: Bad Now, Worse Tomorrow



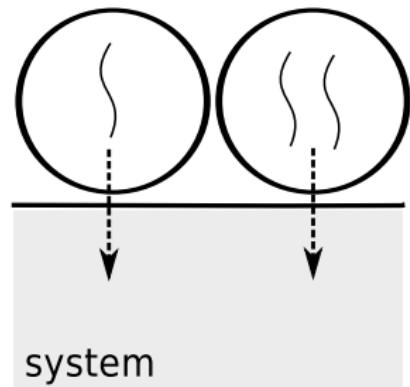
Decreasing process sizes → 5nm

- + faster
- + less power
- + smaller
- increased vulnerability to HW transient faults

# Application-Level Fault Tolerance

## Application fault tolerance

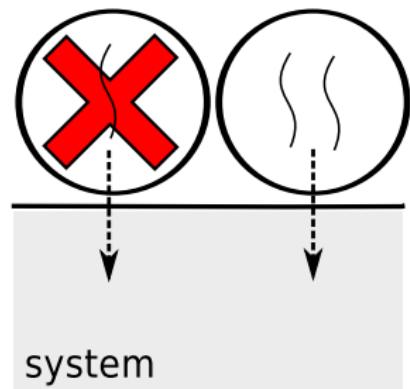
- example recovery techniques
  - recovery blocks
  - checkpointing
  - re-fork
- temporal redundancy
  - detect fault by job completion
  - replay execution from saved state
- re-execution impacts only lower-priority tasks



# Application-Level Fault Tolerance

## Application fault tolerance

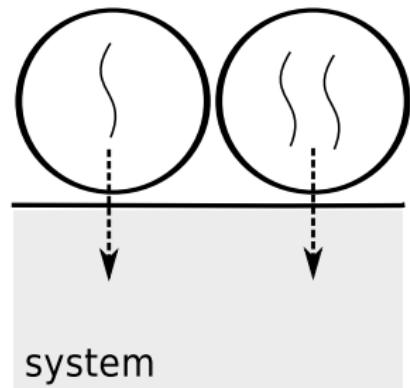
- example recovery techniques
  - recovery blocks
  - checkpointing
  - re-fork
- temporal redundancy
  - detect fault by job completion
  - replay execution from saved state
- re-execution impacts only lower-priority tasks



# Application-Level Fault Tolerance

## Application fault tolerance

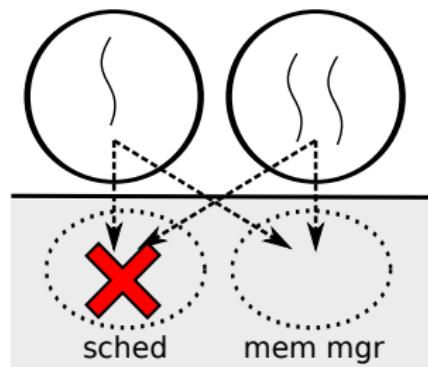
- example recovery techniques
  - recovery blocks
  - checkpointing
  - re-fork
- temporal redundancy
  - detect fault by job completion
  - replay execution from saved state
- re-execution impacts only lower-priority tasks



# System-Level Fault Tolerance

## System-level fault tolerance

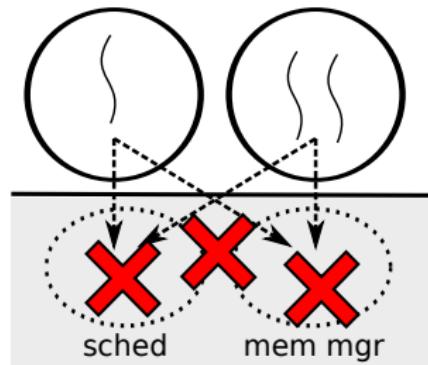
- failures in
  - scheduler
  - memory mapping manager
  - file-systems
  - synchronization manager
  - ...



# System-Level Fault Tolerance

## System-level fault tolerance

- failures in
  - scheduler
  - memory mapping manager
  - file-systems
  - synchronization manager
  - ...



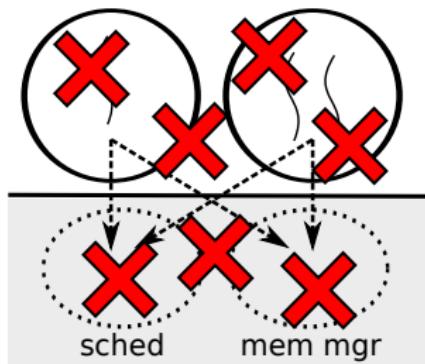
System components contain state for all tasks

- failure impacts memory of *all* tasks

# System-Level Fault Tolerance

## System-level fault tolerance

- failures in
  - scheduler
  - memory mapping manager
  - file-systems
  - synchronization manager
  - ...



System components contain state for all tasks

- failure impacts memory of *all* tasks

Recovery requires resources

- processing time...to recover scheduler
- memory...to recover memory mapper

# System-Level Fault Tolerance

## System-level fault tolerance

- failures in
  - scheduler
  - memory mapping manager
  - file-systems
  - synchronization manager
  - ...

System components contain state for all tasks

- failure impacts memory of *all* tasks



Recovery requires resources

- processing time...to recover scheduler
- memory...to recover memory mapper

# $C^3$ : The Computational Crash Cart

## $C^3$ : Computational Crash Cart

- resuscitate system
- from system-level faults
- predictably



### Main ideas

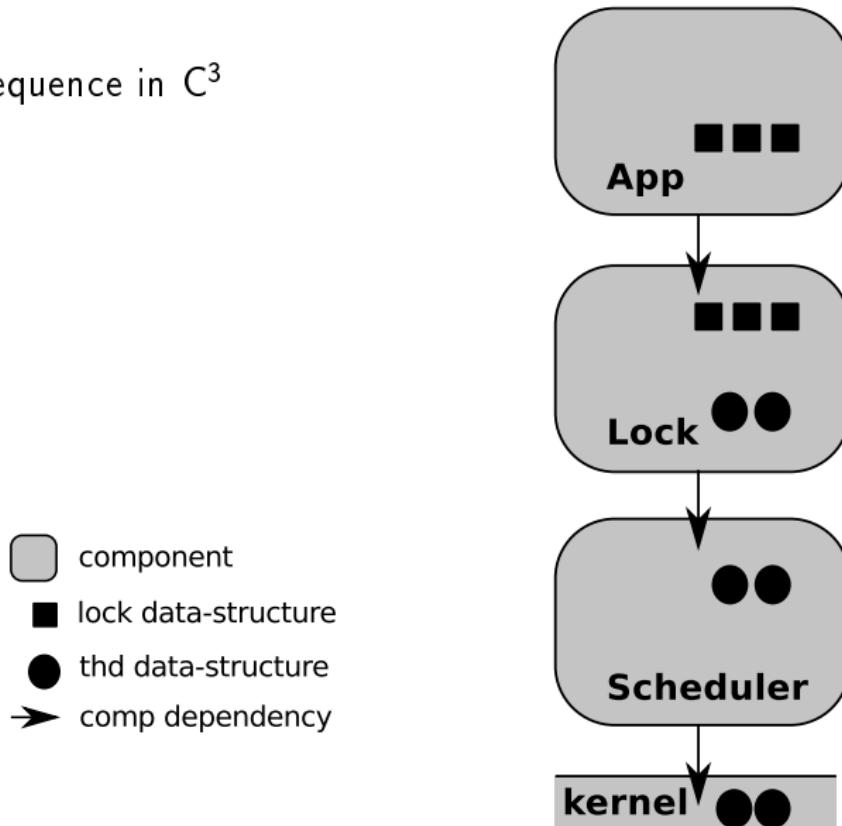
- pervasive **fault isolation** → restrict propagation
- efficient  **$\mu$ -reboot** of individual components
- **interface-driven**, application-oblivious recovery
- on-demand recovery → bound **recovery inversion**

$C^3$  – an interface-driven, predictable system-level fault recovery mechanism

J. Song, J. Wittrock, and G. Palmer, “Predictable, efficient system-level fault tolerance in  $C^3$ ” in RTSS, 2013

# $C^3$ System-Level Fault Recovery

Recovery sequence in  $C^3$

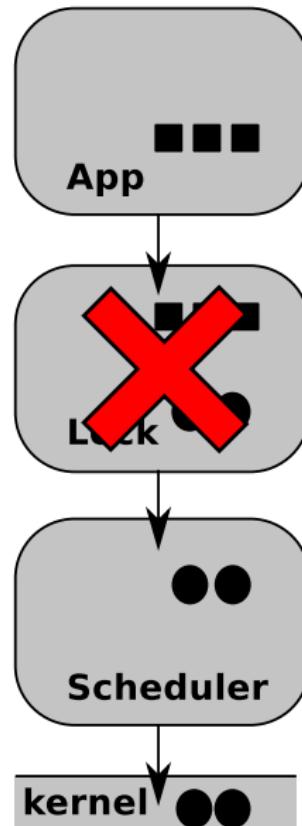


# $C^3$ System-Level Fault Recovery

Recovery sequence in  $C^3$

- 1 fault detection

- component
- lock data-structure
- thd data-structure
- comp dependency

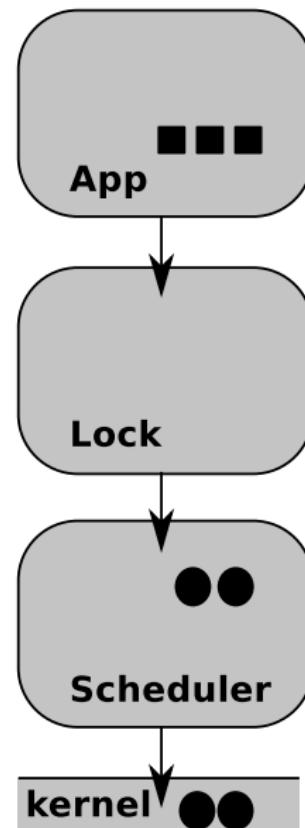


# $C^3$ System-Level Fault Recovery

Recovery sequence in  $C^3$

- 1 fault detection
- 2 safe-state recovery
  - $\mu$ -reboot

- component
- lock data-structure
- thd data-structure
- comp dependency

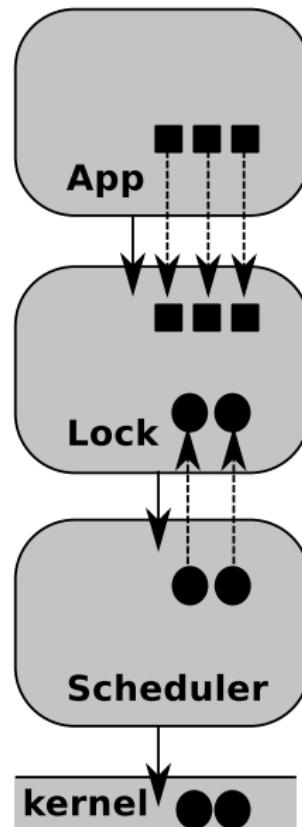


# $C^3$ System-Level Fault Recovery

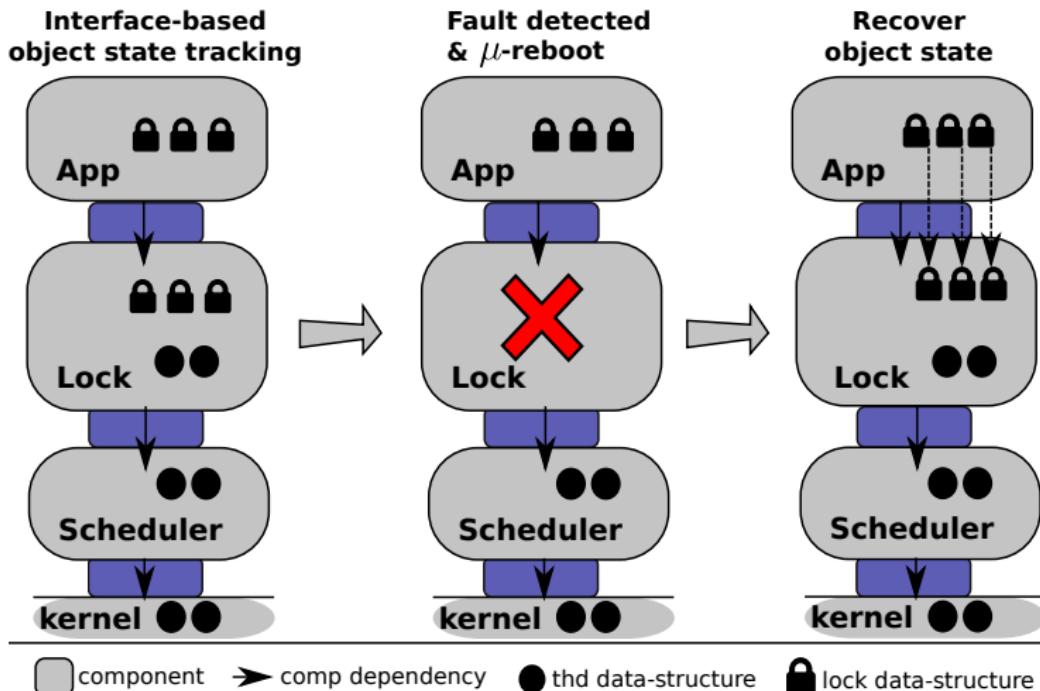
Recovery sequence in  $C^3$

- 1 fault detection
- 2 safe-state recovery
  - $\mu$ -reboot
- 3 consistent-state recovery
  - object state recovery

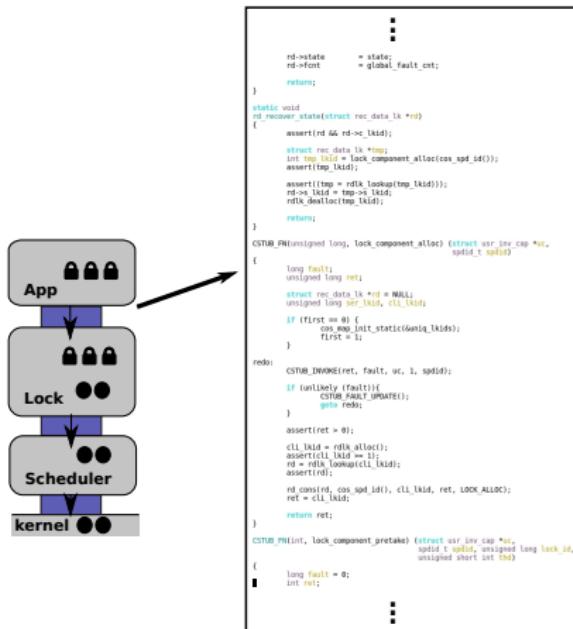
- component
- lock data-structure
- thd data-structure
- comp dependency



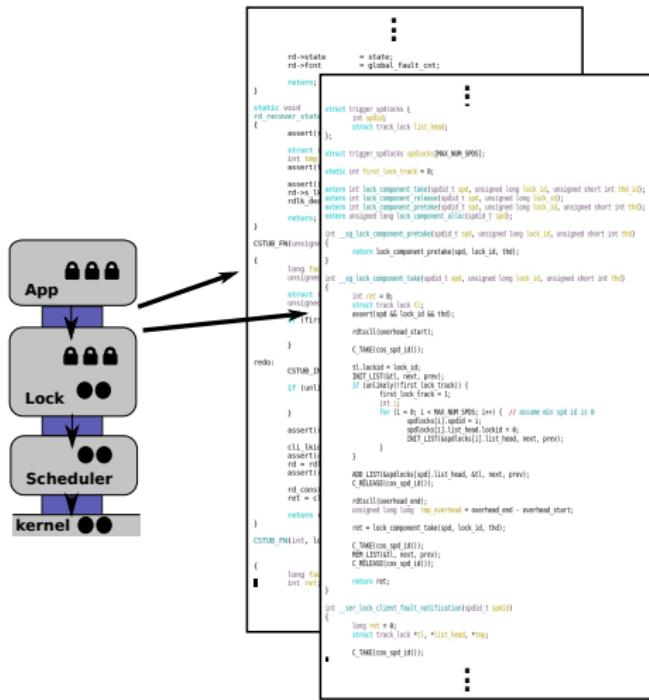
# C<sup>3</sup> System-Level Fault Recovery



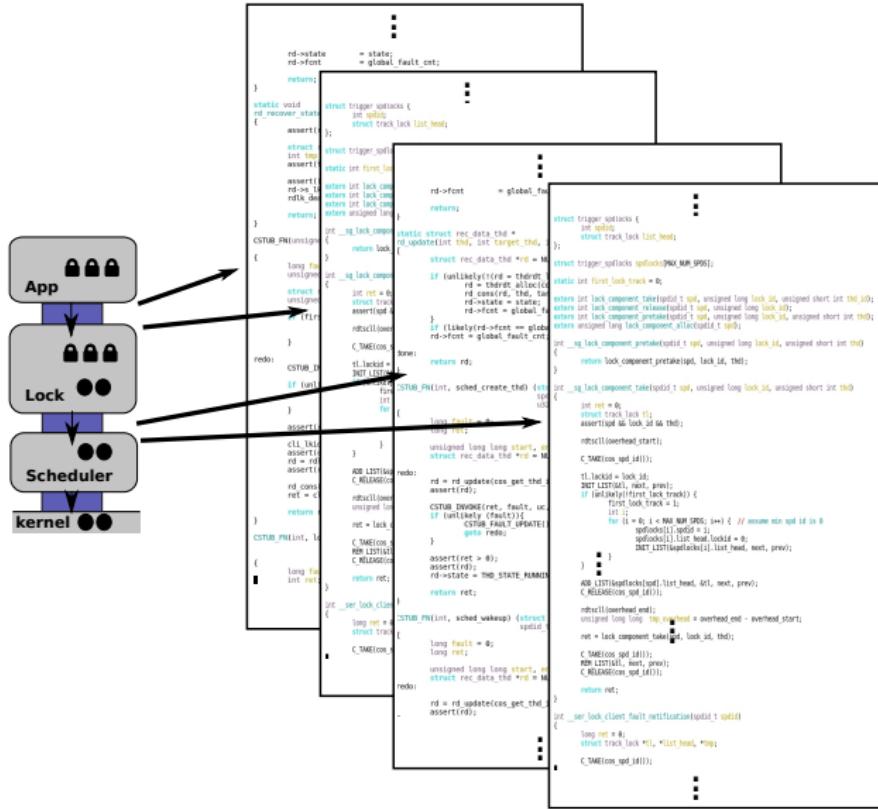
# C<sup>3</sup> Implementation – Writing Code Manually



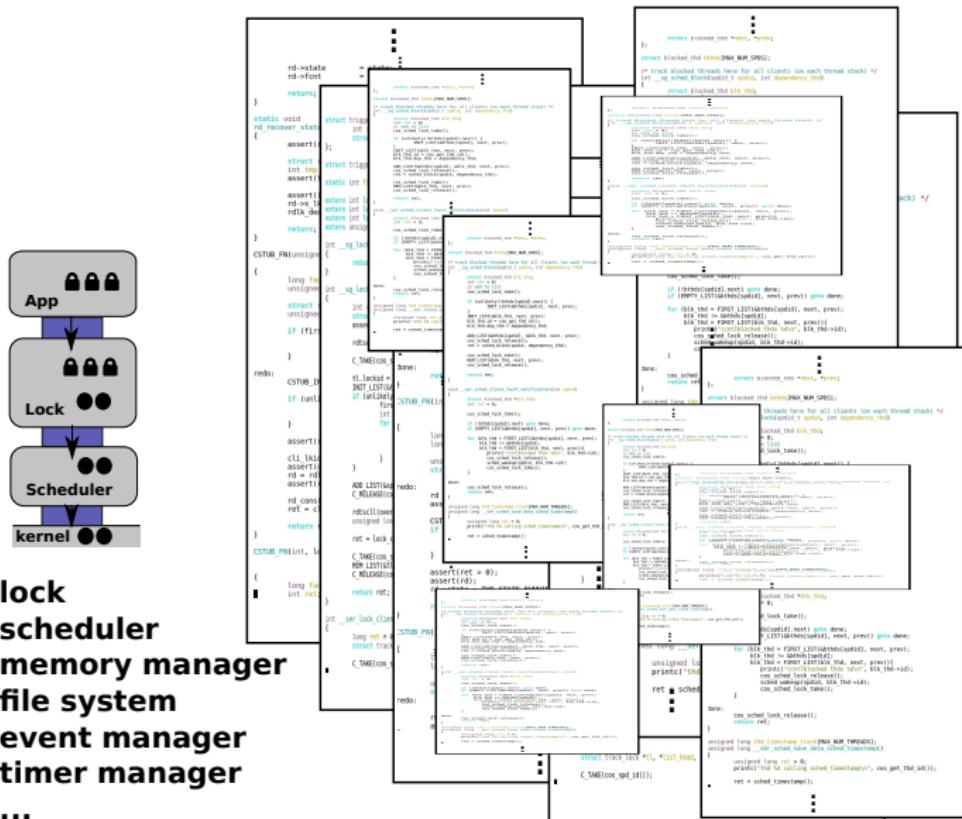
## C<sup>3</sup> Implementation – Writing Code Manually



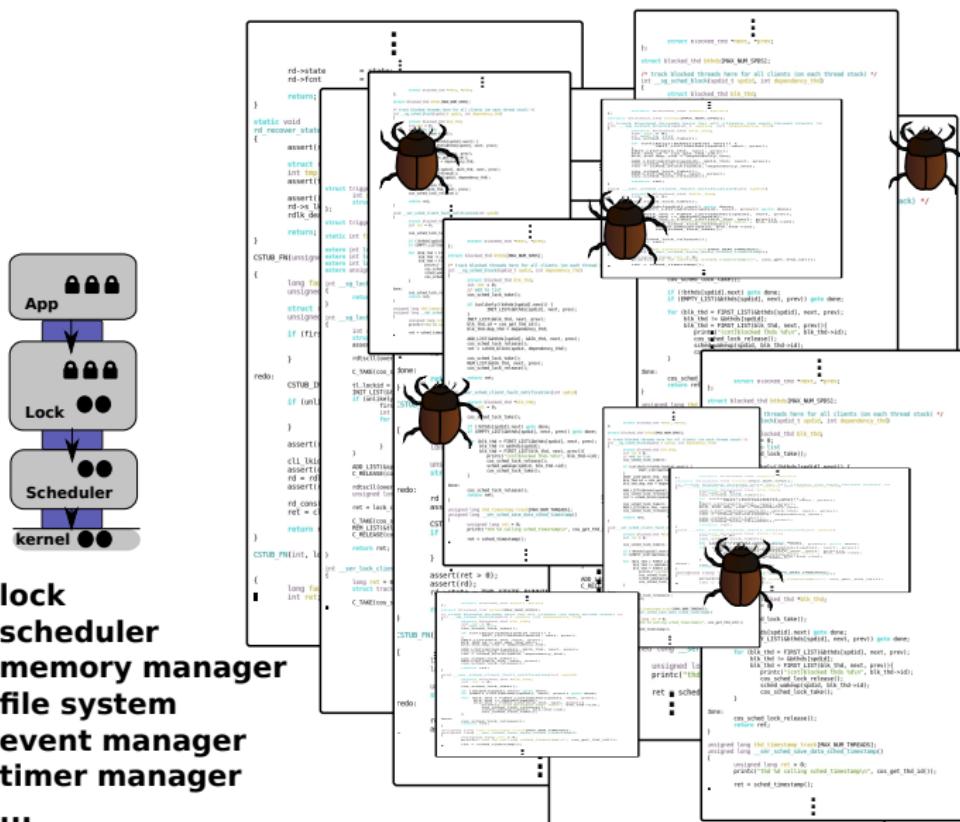
## C<sup>3</sup> Implementation – Writing Code Manually



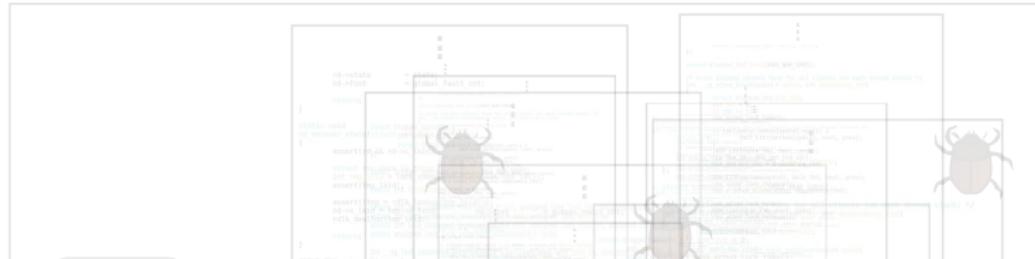
## C<sup>3</sup> Implementation – Writing Code Manually



## C<sup>3</sup> Implementation – Writing Code Manually



# C<sup>3</sup> Implementation – Writing Code Manually



Manually writing is ad-hoc and error-prone

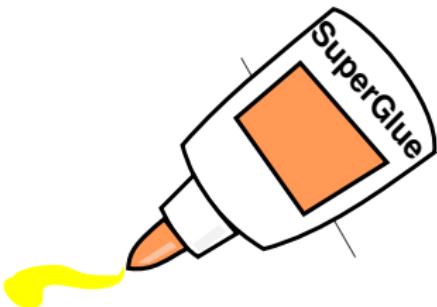
Goal → automatically generate C<sup>3</sup>-style recovery code

lock  
scheduler  
memory manager  
file system  
event manager  
timer manager  
...



## SuperGlue

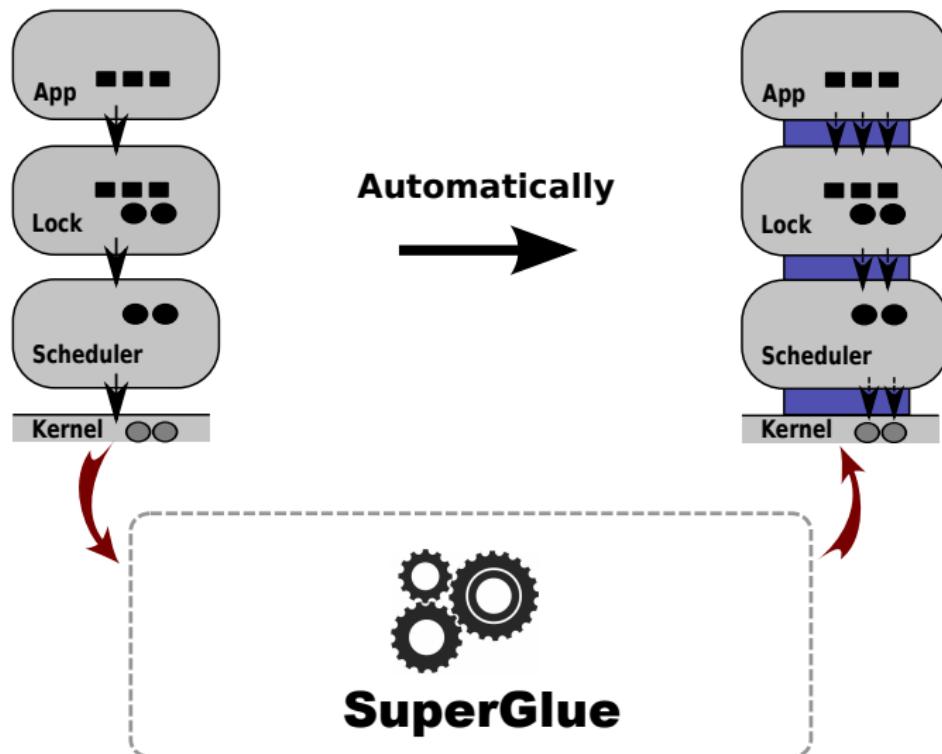
- automatically generate interface-driven recovery code
- for system-level services



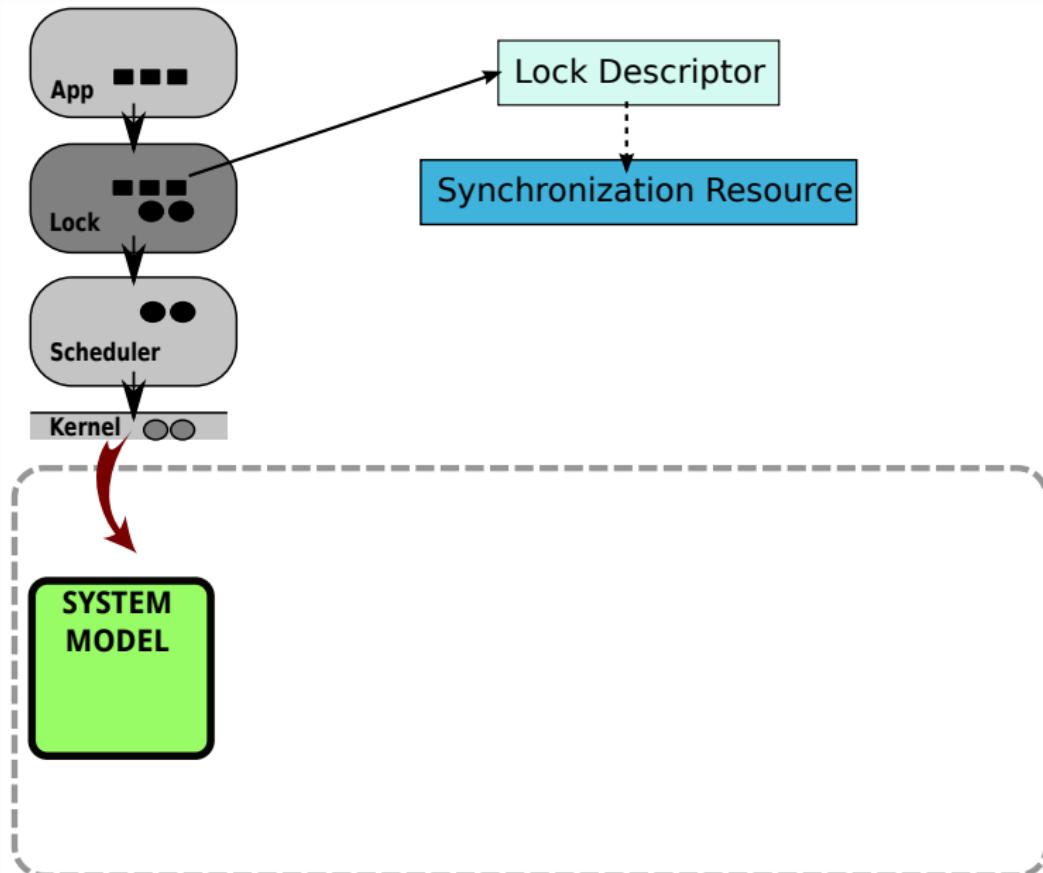
## Main ideas

- A system **model** and **interface description language (IDL)**
- An **IDL compiler** synthesizes recovery code from the model

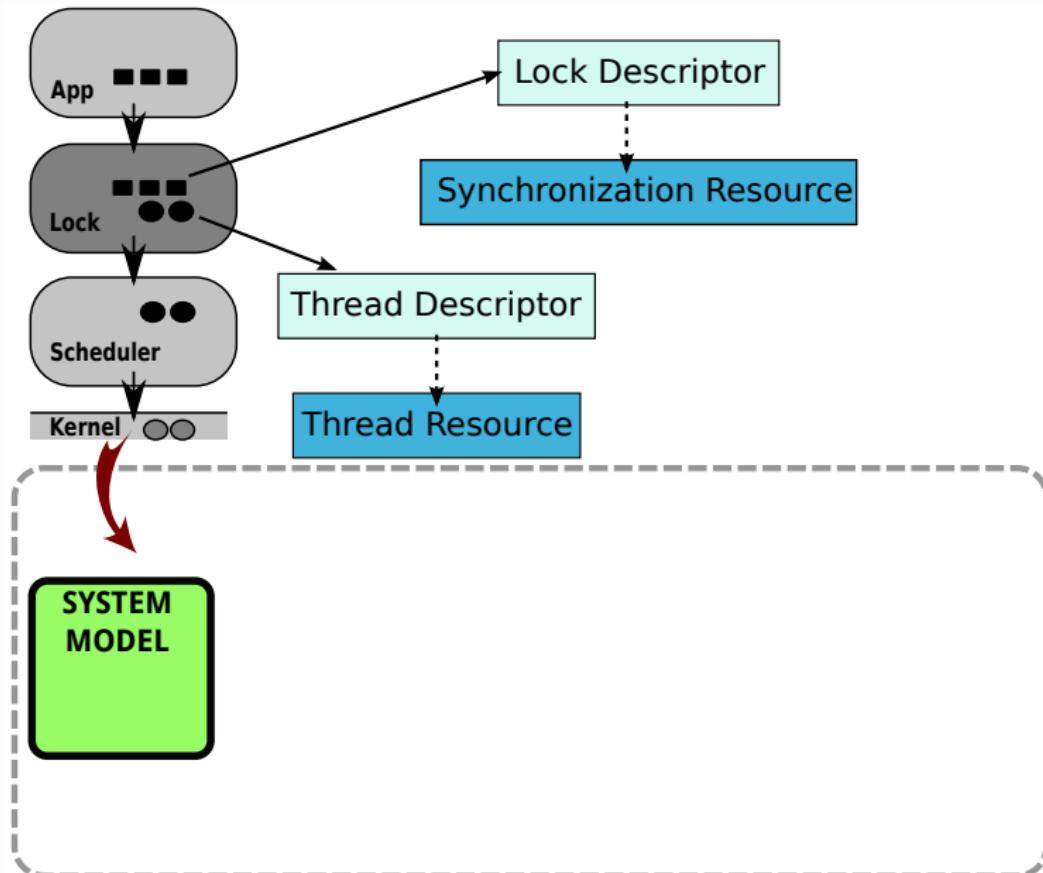
# SuperGlue – IDL-based System-Level Fault Tolerance



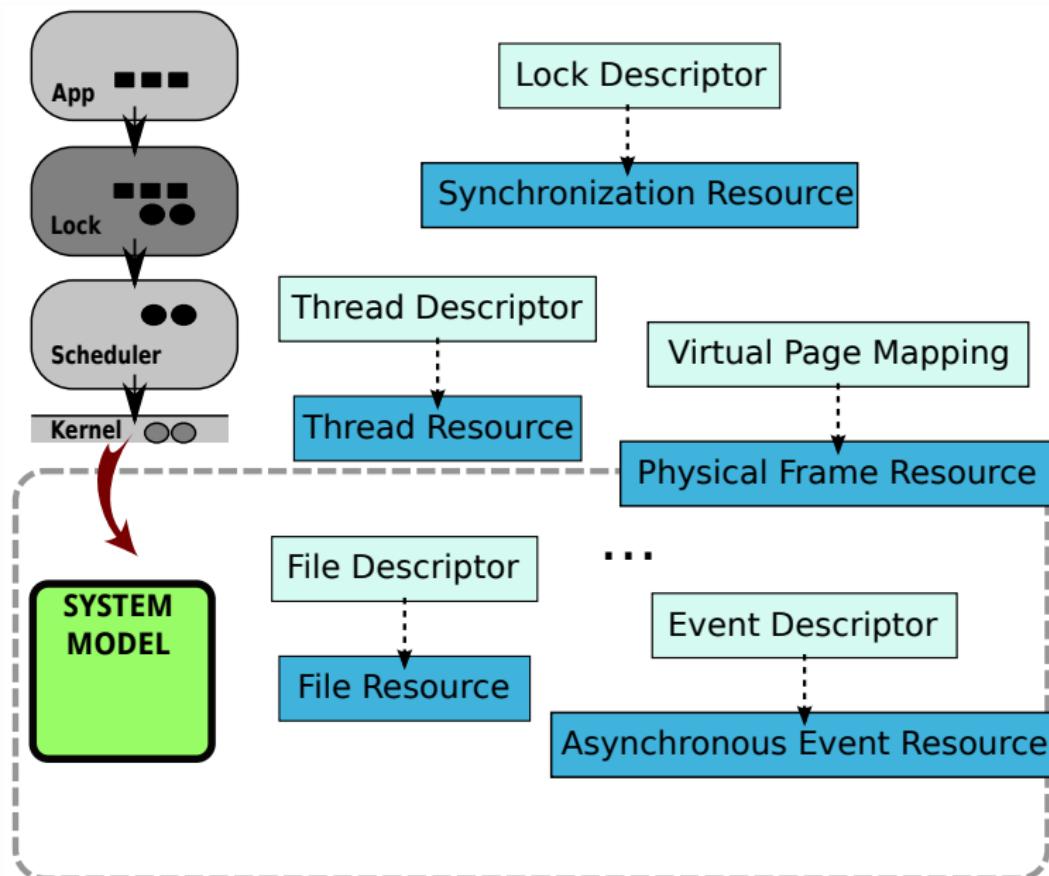
# Descriptor-Resource Model



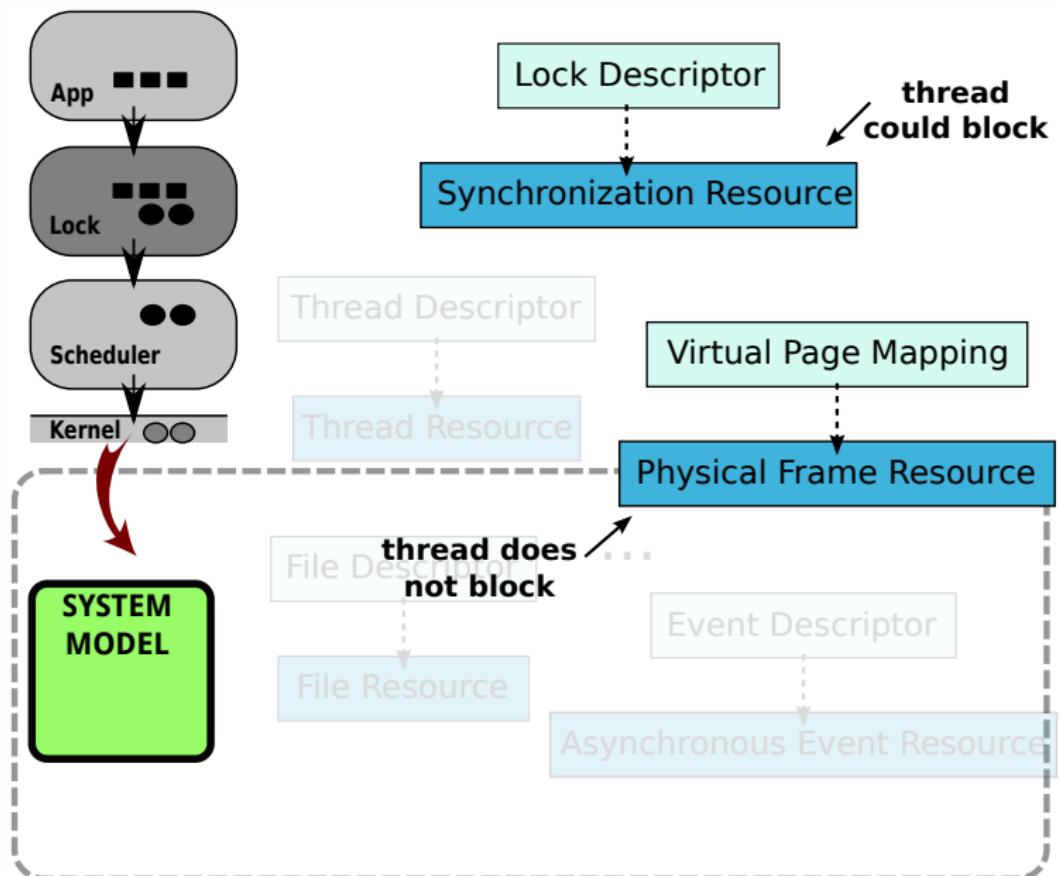
# Descriptor-Resource Model



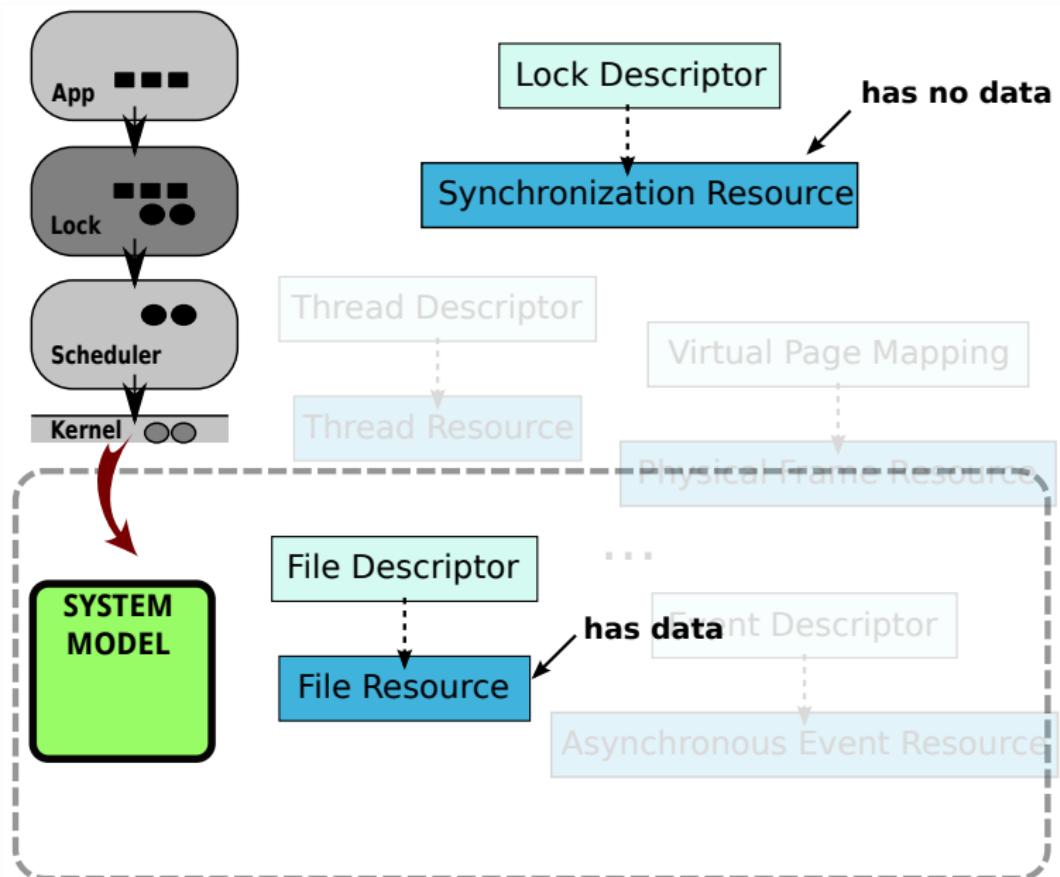
# Descriptor-Resource Model



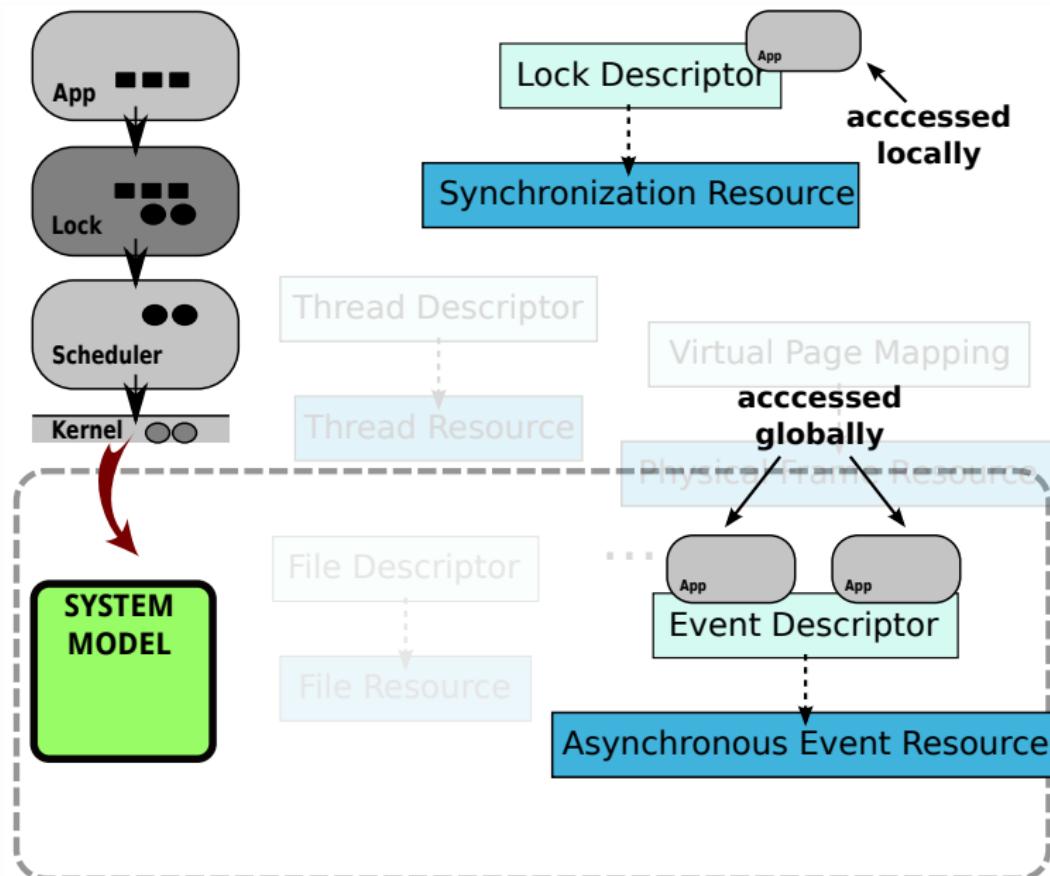
# Descriptor-Resource Model



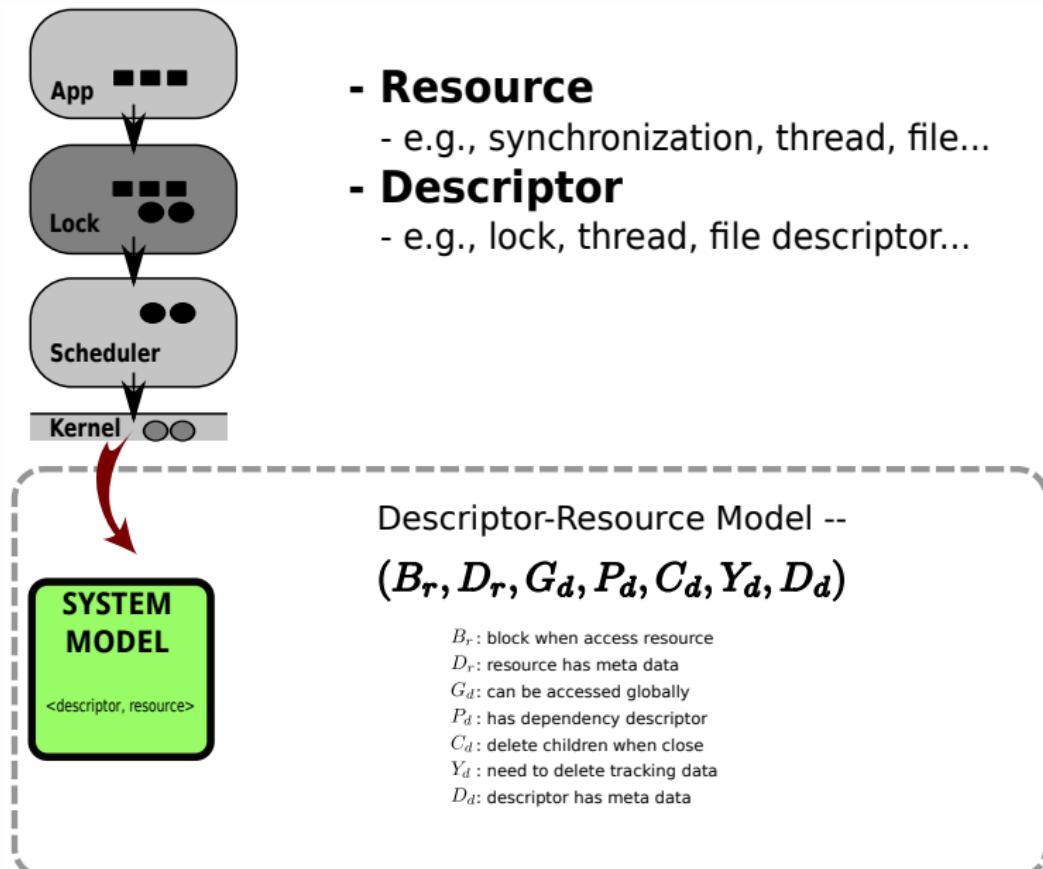
# Descriptor-Resource Model



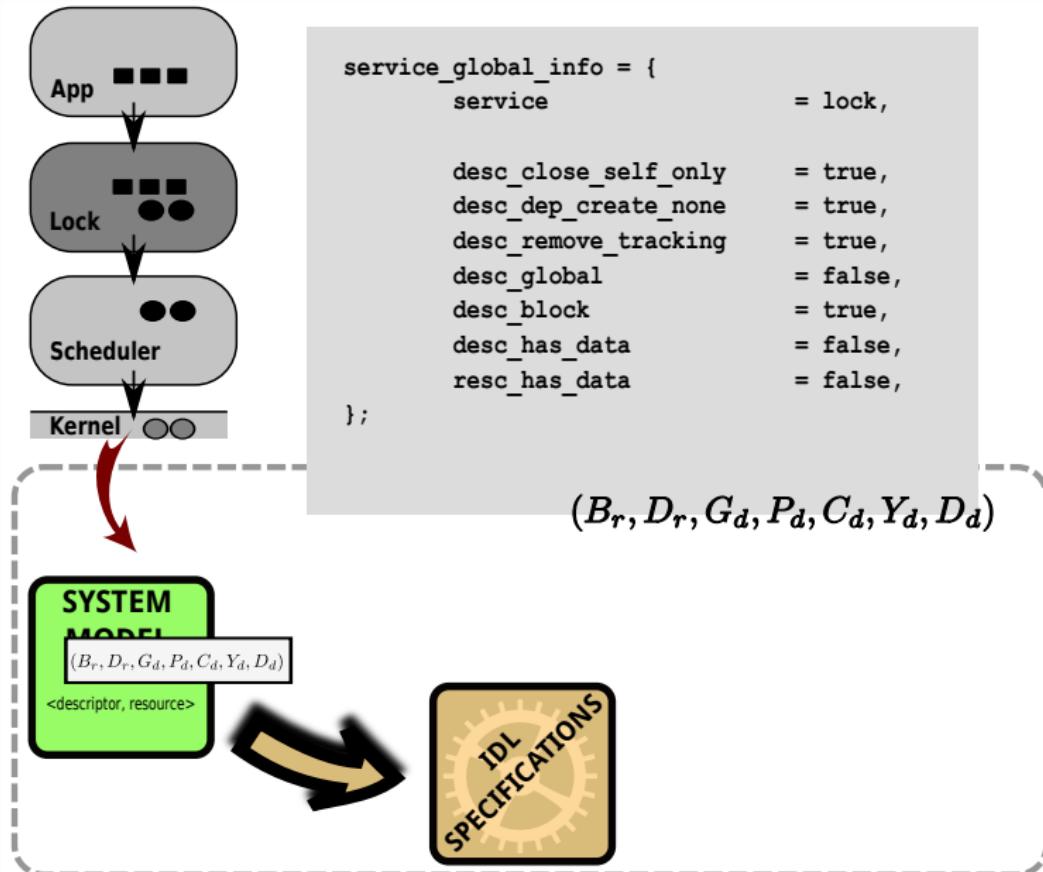
# Descriptor-Resource Model



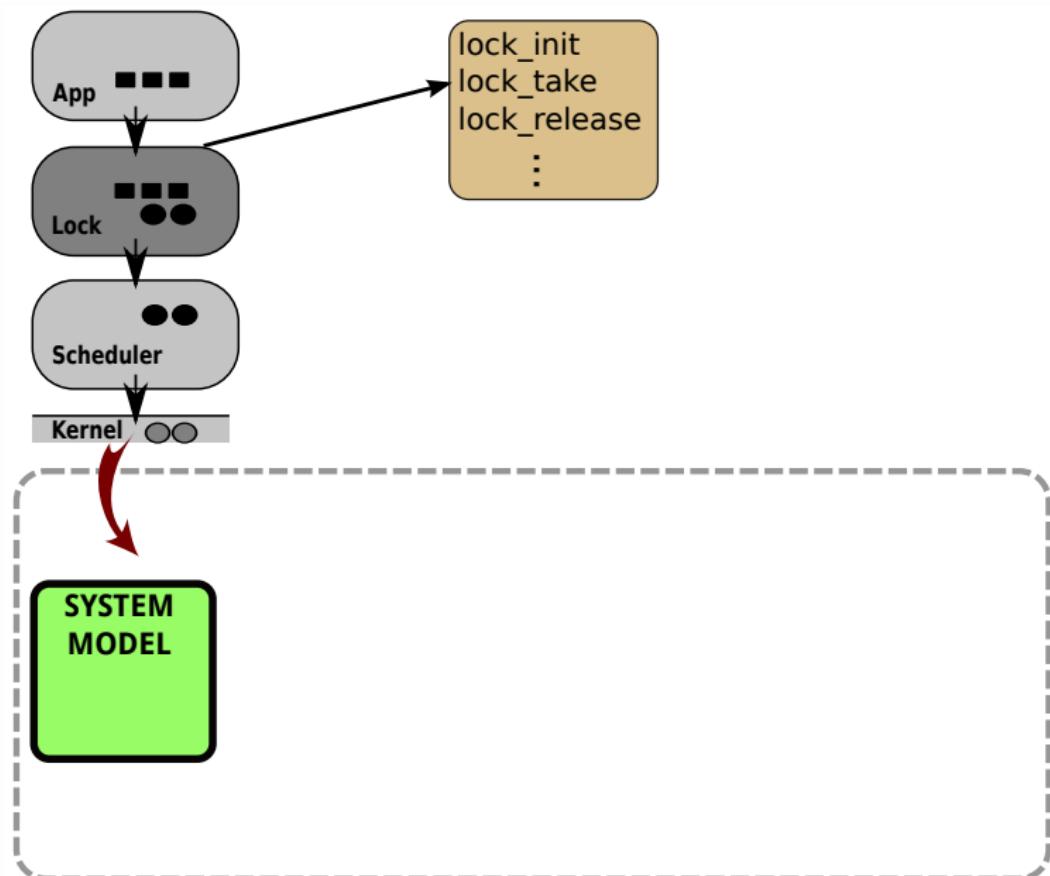
# Descriptor-Resource Model



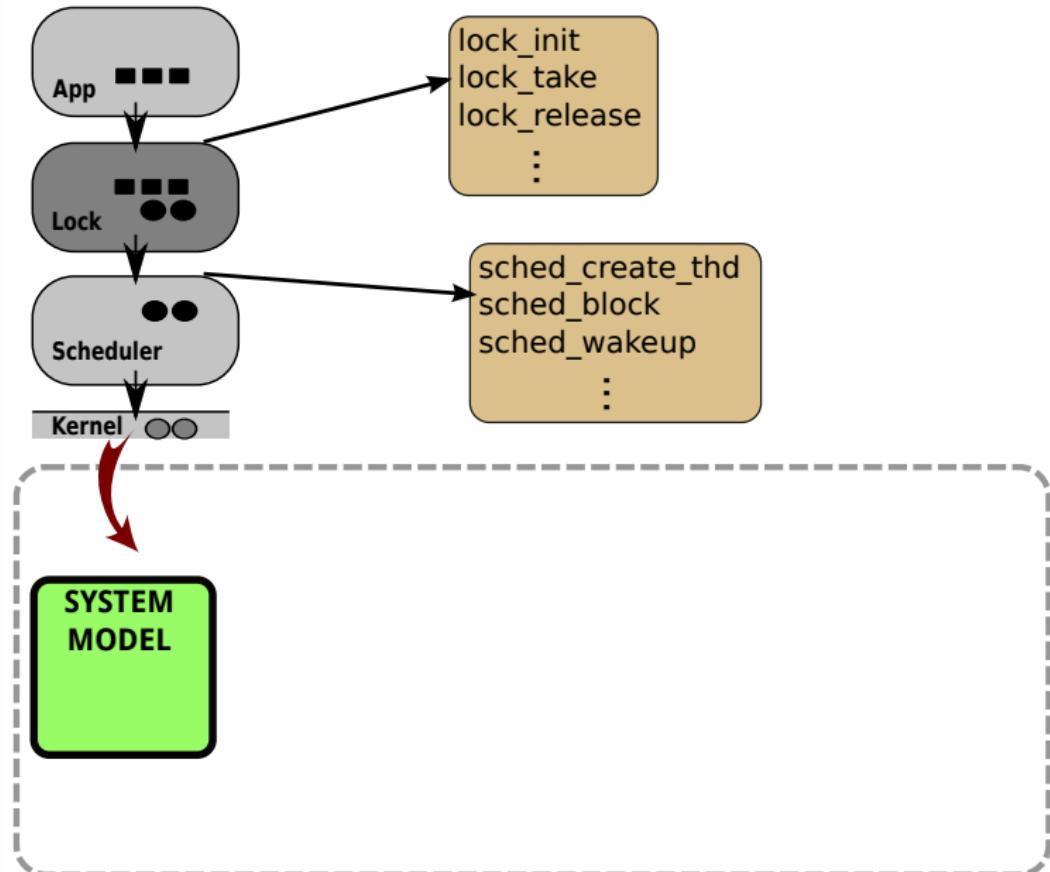
# Descriptor-Resource Model → IDL



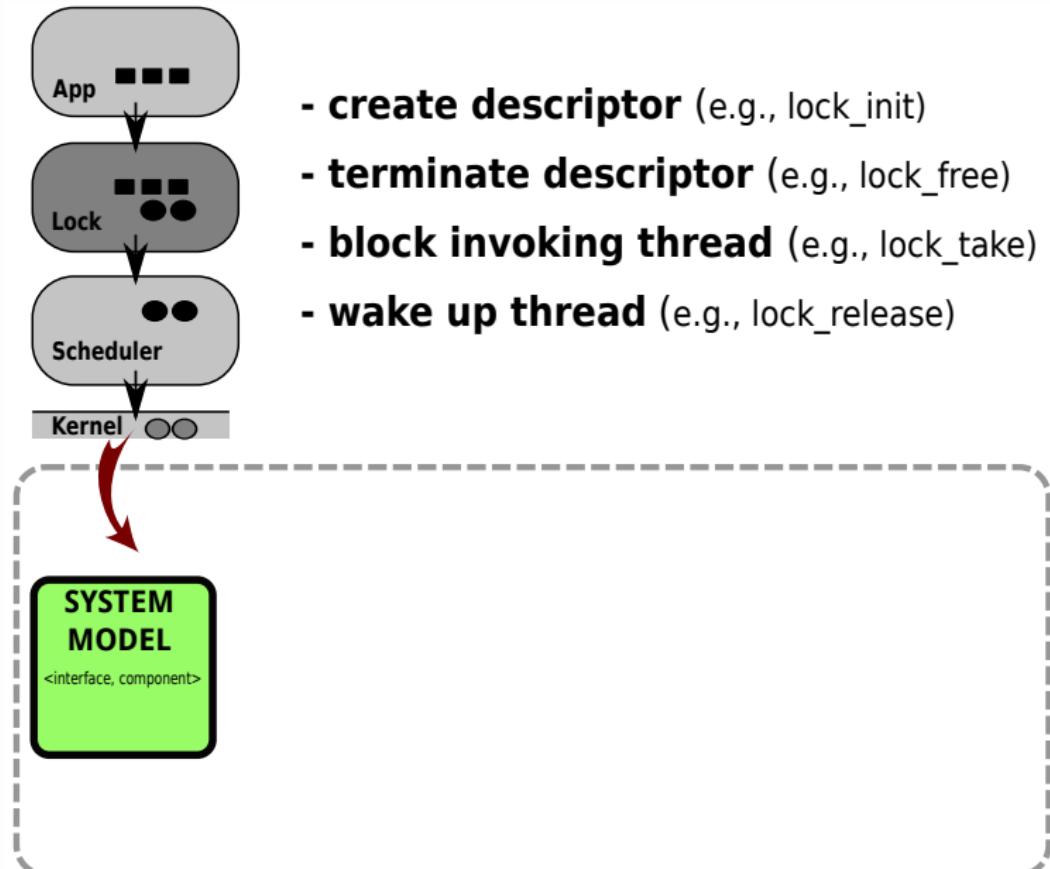
# Component Interface Function



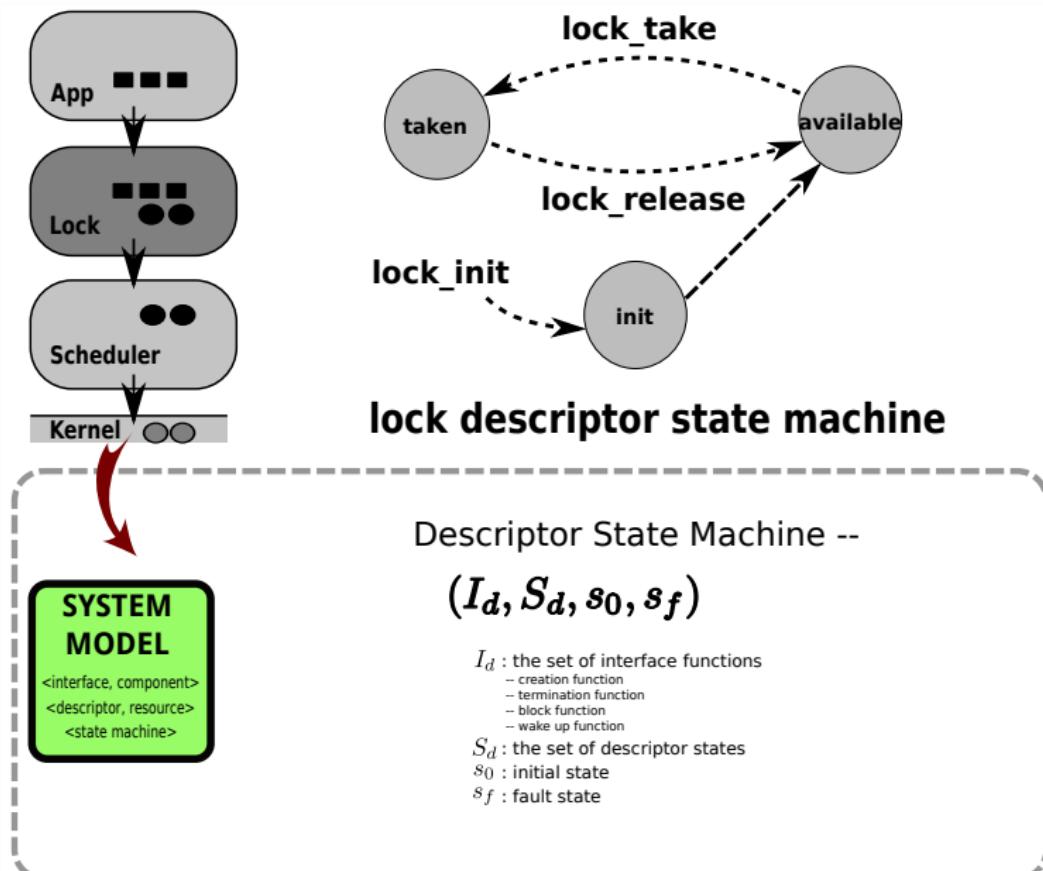
# Component Interface Function



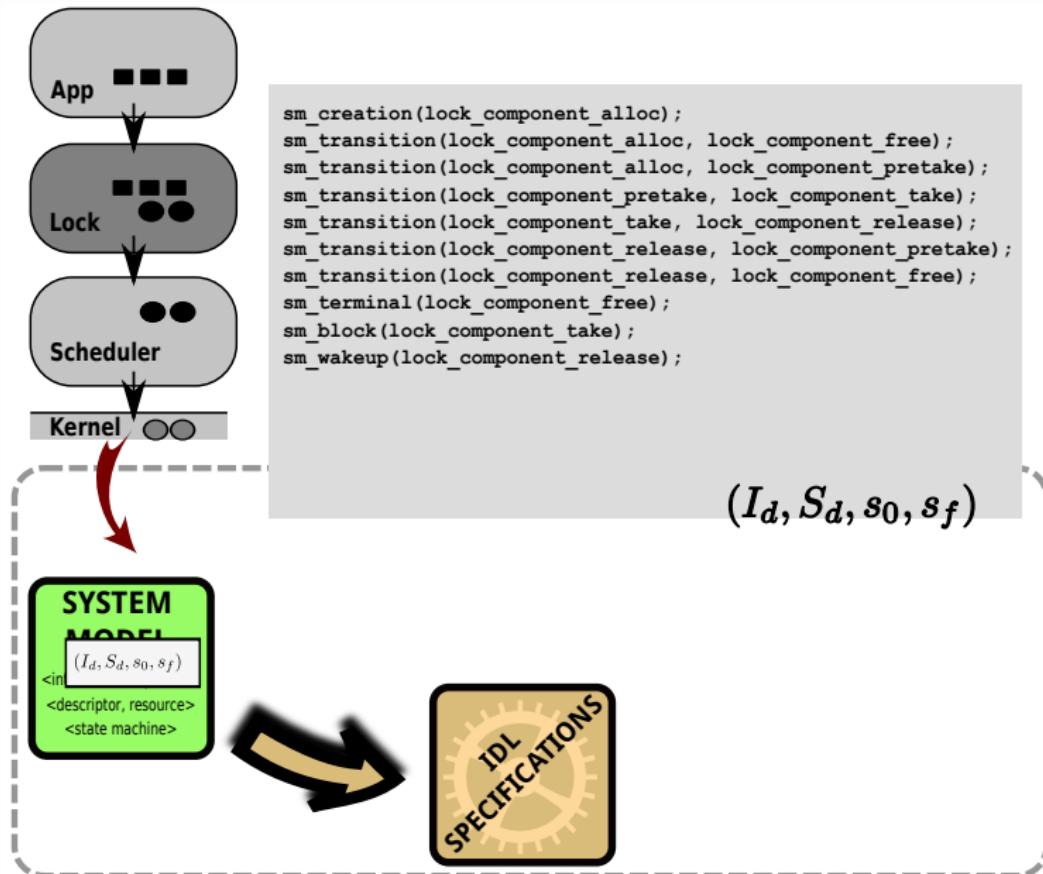
# Component Interface Function



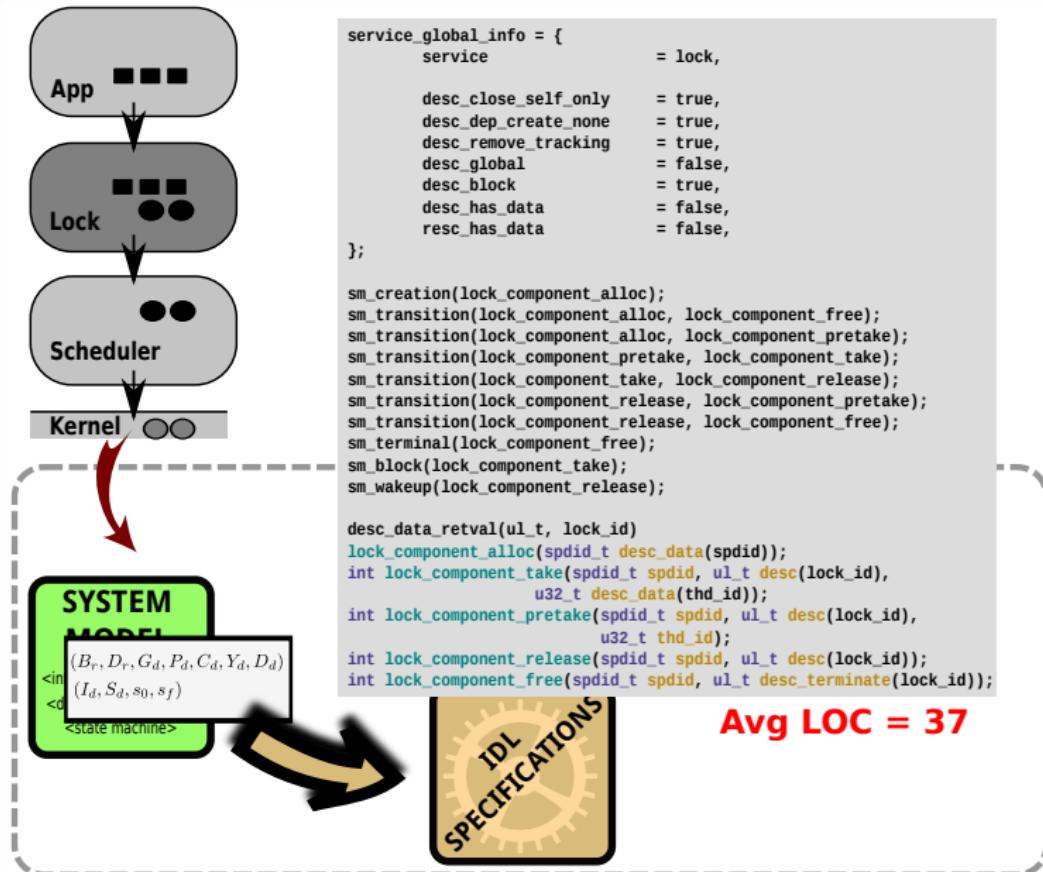
# Descriptor State Machine



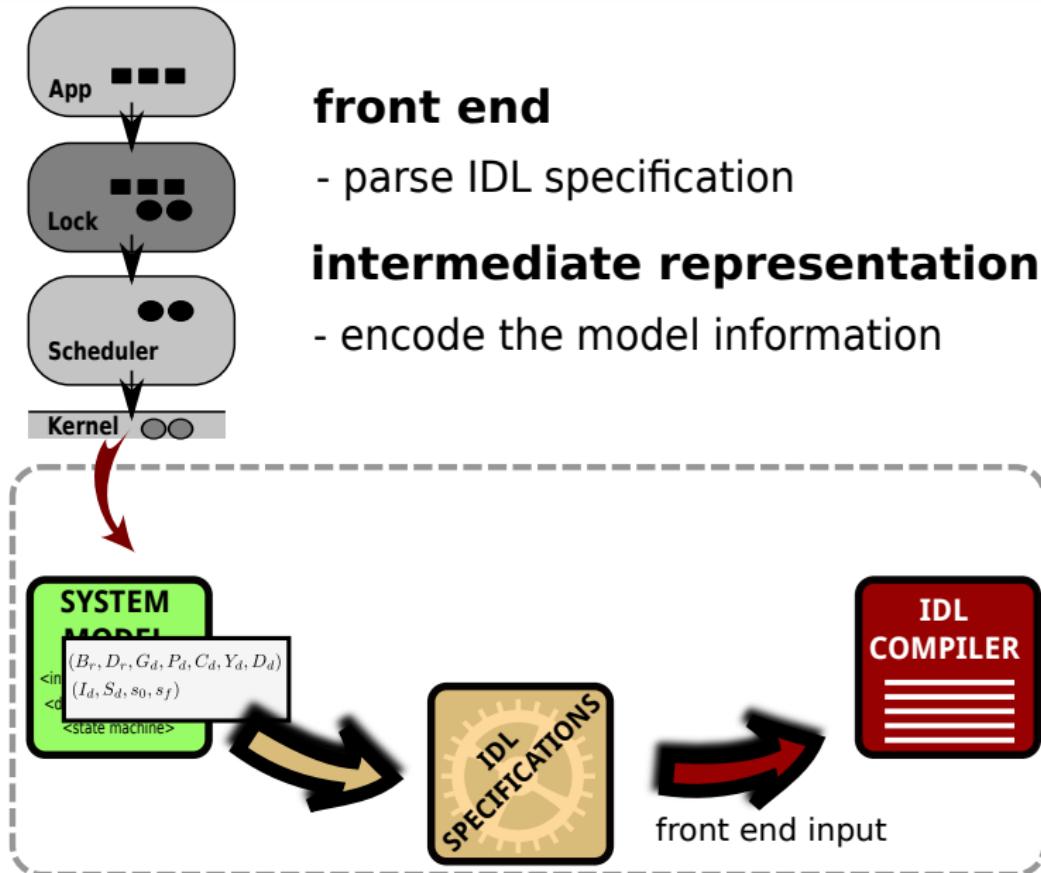
# Descriptor State Machine → IDL



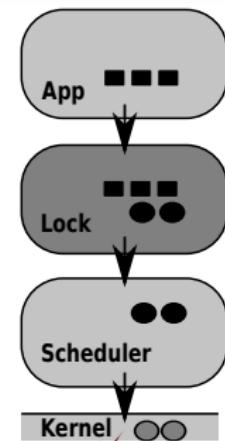
# IDL-based Specification



# Toward Generating Recovery Code

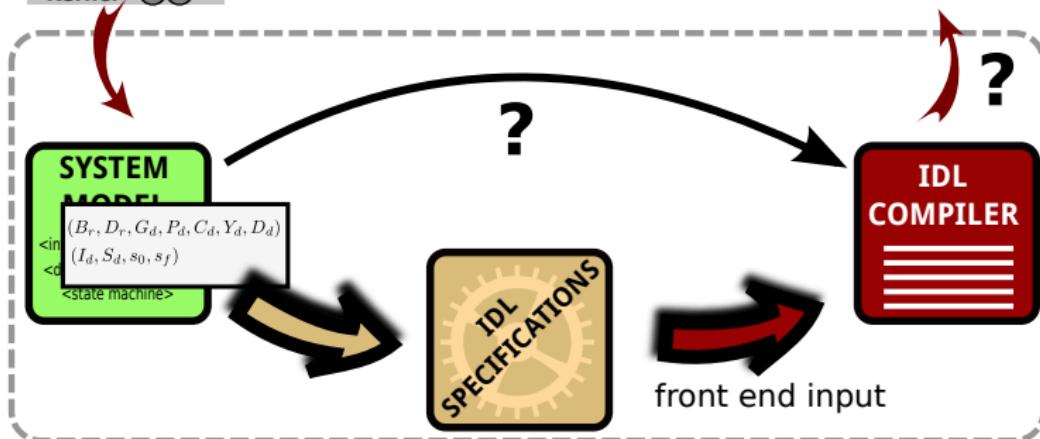


# Toward Generating Recovery Code

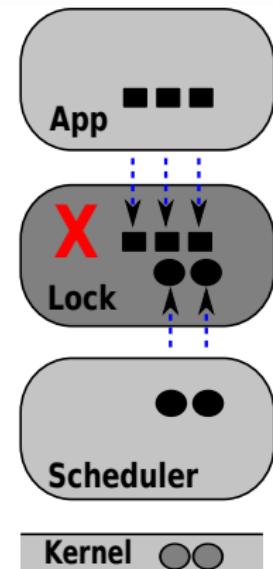


## back end

- what recovery mechanism should be used?
- how to synthesize the code?



# Interface-Driven Recovery Mechanisms

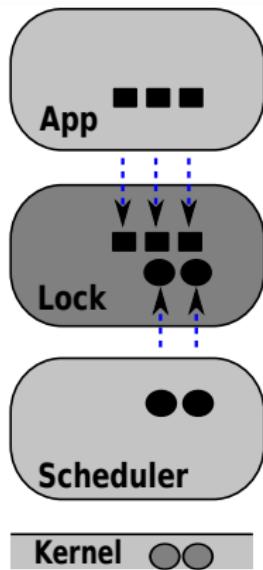


basic recovery

-- through component operation (**always**)

basic  
recovery

# Interface-Driven Recovery Mechanisms



basic recovery

-- through component operation (**always**)

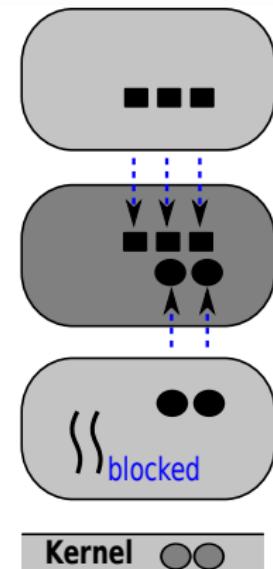
timing of recovery - on-demand

-- recover only when accessed (**always**)

basic  
recovery

timing  
of recovery

# Interface-Driven Recovery Mechanisms



## basic recovery

-- through component operation (**always**)

## timing of recovery - on-demand

-- recover only when accessed (**always**)

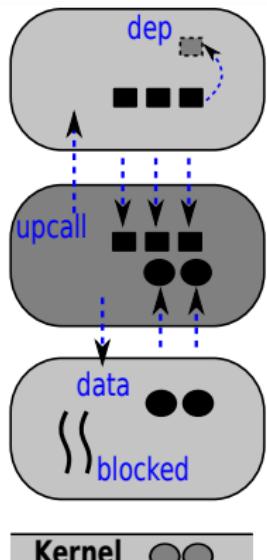
## timing of recovery - eager

-- eagerly wake up blocking threads (**B<sub>r</sub>**)

---

basic recovery	timing of recovery
----------------	--------------------

# Interface-Driven Recovery Mechanisms



## basic recovery

- through component operation (**always**)

## timing of recovery - on-demand

- recover only when accessed (**always**)

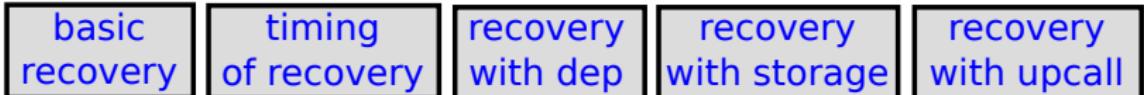
## timing of recovery - eager

- eagerly wake up blocking threads ( $B_r$ )

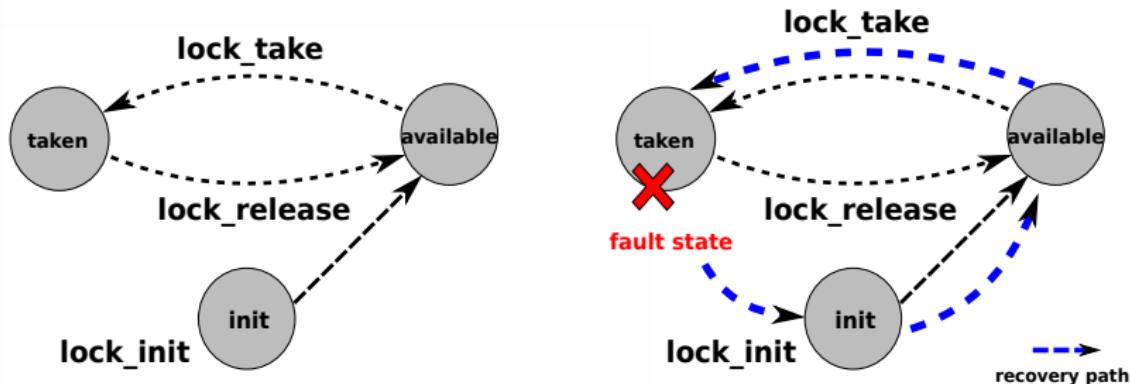
## recovery with dependency

- require to reconstruct parent ( $P_d$ )
- require to reconstruct children ( $C_d$ )

:

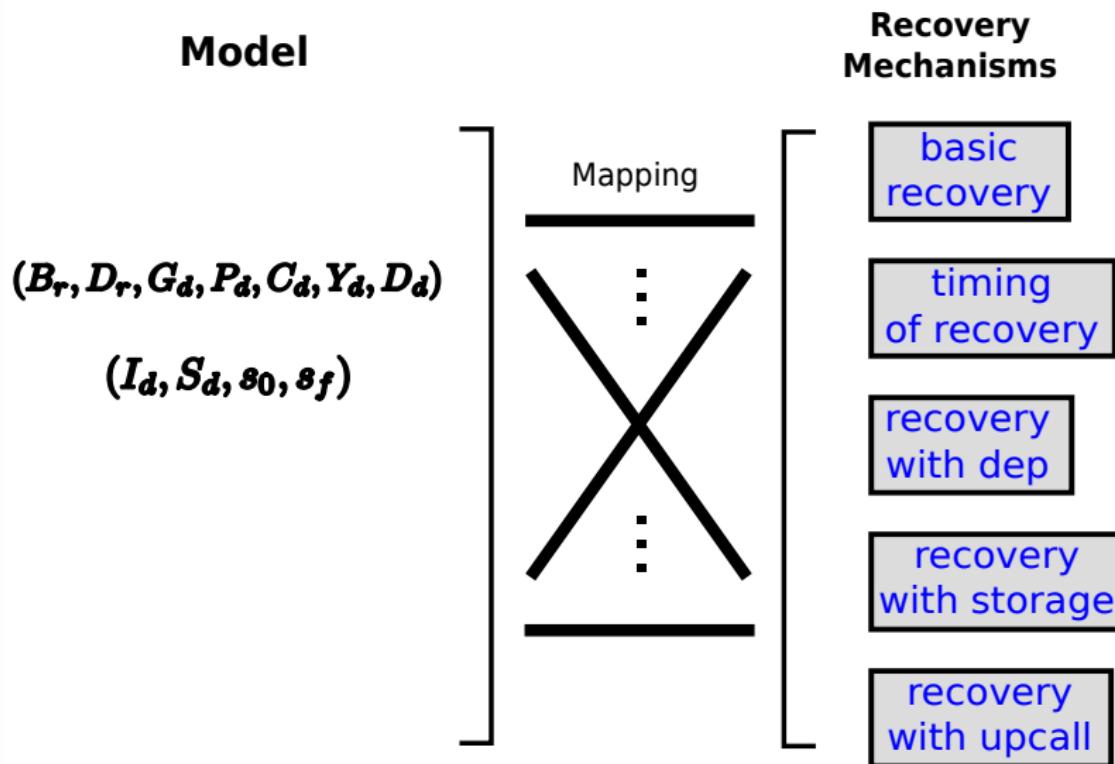


# Interface-Driven Recovery Mechanisms



Descriptor State Machine

# Model → Recovery Mechanisms



# Synthesize the Recovery Code

## Predicates

$\vdots$   
 $true$   
 $create\_fn \wedge \neg G_d$   
 $wakeup\_fn \wedge B_r$   
 $block\_fn \wedge B_r$   
 $terminate\_fn \wedge C_d$   
 $G_d \wedge P_d \wedge C_d \wedge D_d$   
 $D_r \wedge \neg P_d \wedge Y_d$   
 $\vdots$   
 $\vdots$

## Templates

`;; predicate true :  
CREATE_FN([fn], [fn], [fn], [fn], [fn], [fn], [fn]);  
;; fn : string;  
;; fn : string;`

`;; predicate true :  
CREATE_FN([fn], [fn], [fn], [fn], [fn], [fn], [fn]);  
;; fn : string;  
;; fn : string;`

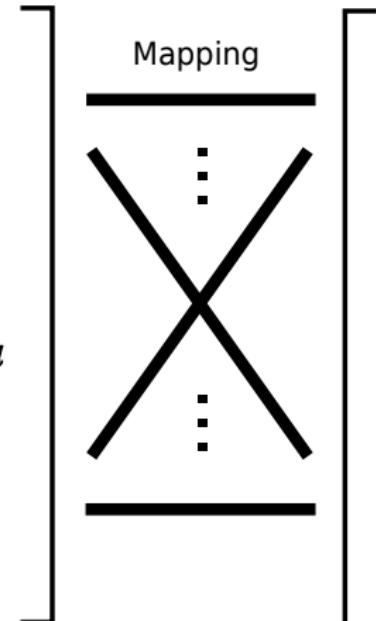
`;; predicate true :  
CREATE_FN([fn], [fn], [fn], [fn], [fn], [fn], [fn]);  
;; fn : string;  
;; fn : string;`

`;; predicate true :  
CREATE_FN([fn], [fn], [fn], [fn], [fn], [fn], [fn]);  
;; fn : string;  
;; fn : string;`

`;; predicate true :  
CREATE_FN([fn], [fn], [fn], [fn], [fn], [fn], [fn]);  
;; fn : string;  
;; fn : string;`

`;; predicate true :  
CREATE_FN([fn], [fn], [fn], [fn], [fn], [fn], [fn]);  
;; fn : string;  
;; fn : string;`

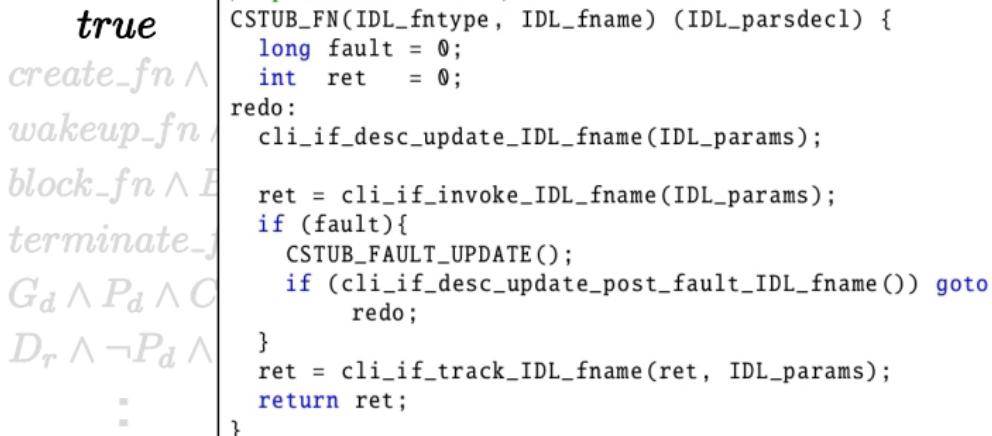
Mapping



# Synthesize the Recovery Code

## Predicates

## Templates



```
/* predicate: true */
CSTUB_FN(IDL_fntype, IDL_fname) (IDL_parsdecl) {
    long fault = 0;
    int ret = 0;
    redo:
        cli_if_desc_update_IDL_fname(IDL_params);

        ret = cli_if_invoke_IDL_fname(IDL_params);
        if (fault){
            CSTUB_FAULT_UPDATE();
            if (cli_if_desc_update_post_fault_IDL_fname()) goto redo;
        }
        ret = cli_if_track_IDL_fname(ret, IDL_params);
}
```

```
/* predicate: true */
CSTUB_FN(IDL_fntype, IDL_fname) (IDL_parsdecl) {
    long fault = 0;
    int ret = 0;
    redo:
        cli_if_desc_update_IDL_fname(IDL_params);

        ret = cli_if_invoke_IDL_fname(IDL_params);
        if (fault){
            CSTUB_FAULT_UPDATE();
            if (cli_if_desc_update_post_fault_IDL_fname()) goto redo;
        }
        ret = cli_if_track_IDL_fname(ret, IDL_params);
}
```

# Synthesize the Recovery Code

## Predicates

*true*

*create\_fn*  $\wedge$

*wakeup\_fn*

*block\_fn*  $\wedge$  *I*

*terminate\_fn*

*G<sub>d</sub>*  $\wedge$  *P<sub>d</sub>*  $\wedge$  *C*

*D<sub>r</sub>*  $\wedge$   $\neg$ *P<sub>d</sub>*  $\wedge$

$\vdots$

$\vdots$

$\vdots$

## Mapping

```
/* predicate: true */
CSTUB_FN(IDL_fntype, IDL_fname) (IDL_parsdecl) {
    long fault = 0;
    int ret = 0;
    redo:
        cli_if_desc_update_IDL_fname(IDL_params);

        ret = cli_if_invoke_IDL_fname(IDL_params);
        if (fault){
            CSTUBFAULT_UPDATE();
            if (cli_if_desc_update_post_fault_IDL_fname()) goto redo;
        }
        ret = cli_if_track_IDL_fname(ret, IDL_params);
        return ret;
}
```

## Templates

```
IDL_fntype( IDL_fntype, IDL_fname, IDL_params);
{
    long fault = 0;
    int ret = 0;
    redo:
        if (fault) {
            if (CSTUBFAULT_UPDATE());
            if (cli_if_desc_update_post_fault_IDL_fname()) goto redo;
        }
        ret = cli_if_track_IDL_fname(ret, IDL_params);
}
```

```
IDL_fntype( IDL_fntype, IDL_fname, IDL_params);
{
    long fault = 0;
    int ret = 0;
    redo:
        if (fault) {
            if (CSTUBFAULT_UPDATE());
            if (cli_if_desc_update_post_fault_IDL_fname()) goto redo;
        }
        ret = cli_if_track_IDL_fname(ret, IDL_params);
}
```

# Synthesize the Recovery Code

## Predicates

=  
:  
:  
*true*  
 $create\_fn \wedge \neg G_d$   
 $wakeup\_fn \wedge \neg G_d$   
 $block\_fn \wedge \neg G_d$   
 $terminate \wedge \neg G_d$   
 $G_d \wedge P_d \wedge \neg D_r$   
 $D_r \wedge \neg P_d$

## Mapping

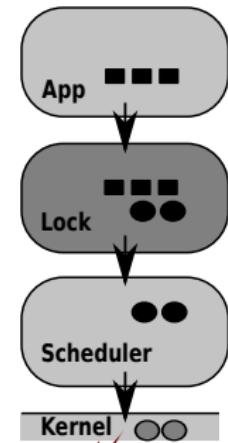
## Templates

```
/* predicate: f ∈ Icreatedr ∧ ¬Gdr */  
static inline int cli_if_track_IDL_fname(int ret,  
                                       IDL_parsdecl) {  
    if (ret == -EINVAL) return ret;  
    struct desc_track *desc = call_desc_alloc();  
    if (!desc) return -ENOMEM;  
    call_desc_track(desc, ret, IDL_params);  
  
    return desc->IDL_id;  
}
```

```
/* predicate: f ∈ Iwakeupdr ∧ ¬Gdr */  
static inline int cli_if_wakeup_IDL_fname(int ret,  
                                         IDL_parsdecl) {  
    if (ret == -EINVAL) return ret;  
    struct desc_track *desc = call_desc_alloc();  
    if (!desc) return -ENOMEM;  
    call_desc_track(desc, ret, IDL_params);  
  
    return desc->IDL_id;
```

```
/* predicate: f ∈ Iblockdr ∧ ¬Gdr */  
static inline int cli_if_block_IDL_fname(int ret,  
                                       IDL_parsdecl) {  
    if (ret == -EINVAL) return ret;  
    struct desc_track *desc = call_desc_alloc();  
    if (!desc) return -ENOMEM;  
    call_desc_track(desc, ret, IDL_params);  
  
    return desc->IDL_id;
```

# Generate Recovery Code



## front end

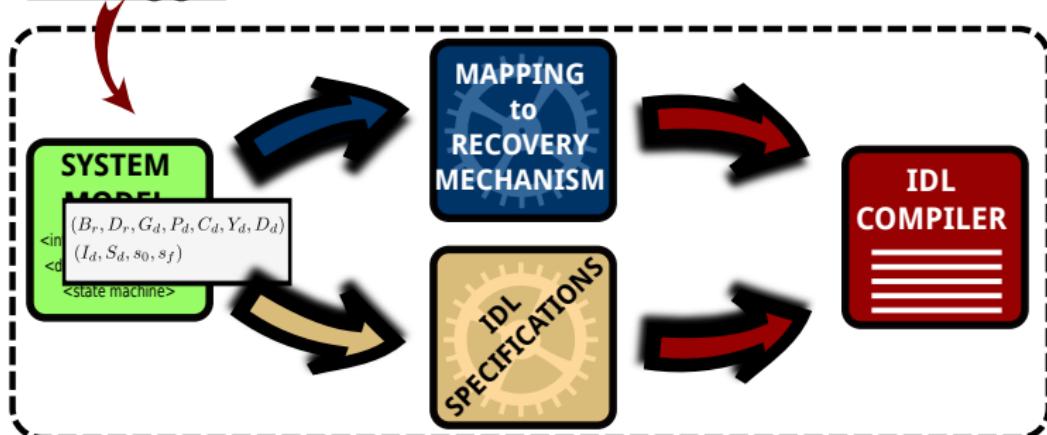
- parse IDL-based specification

## intermediate representation

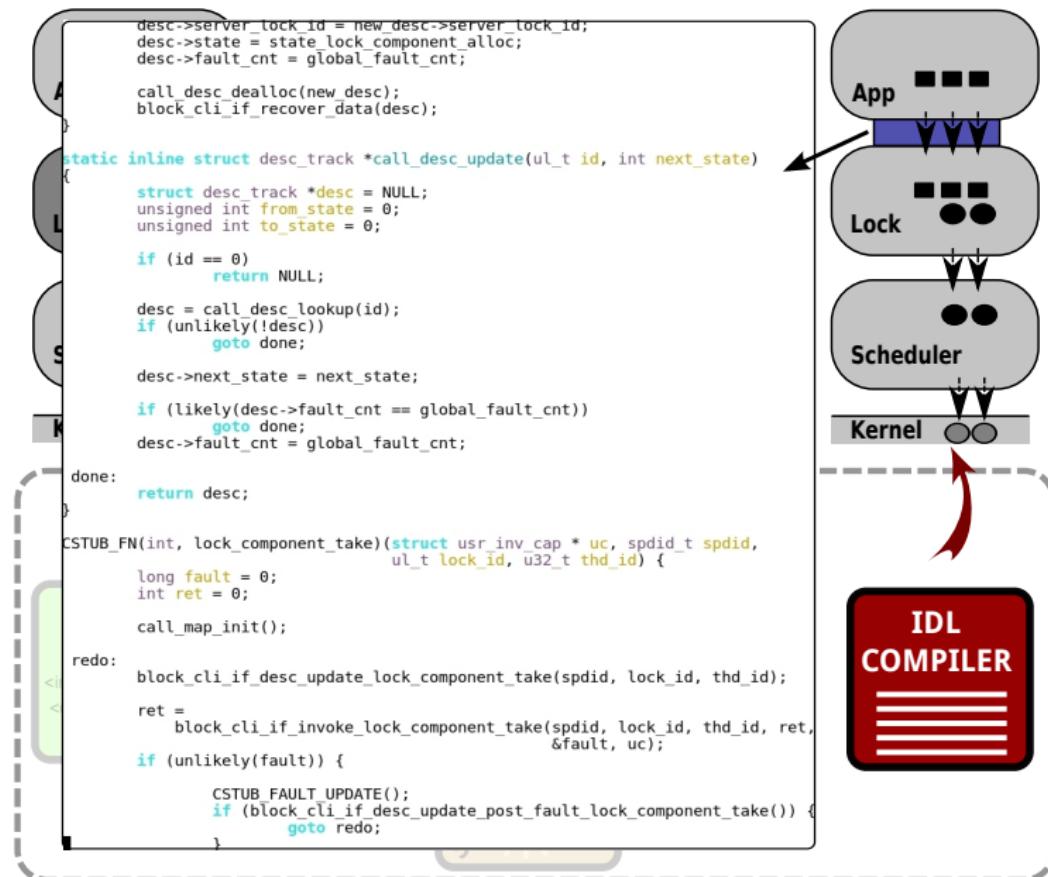
- encode the model information

## back end

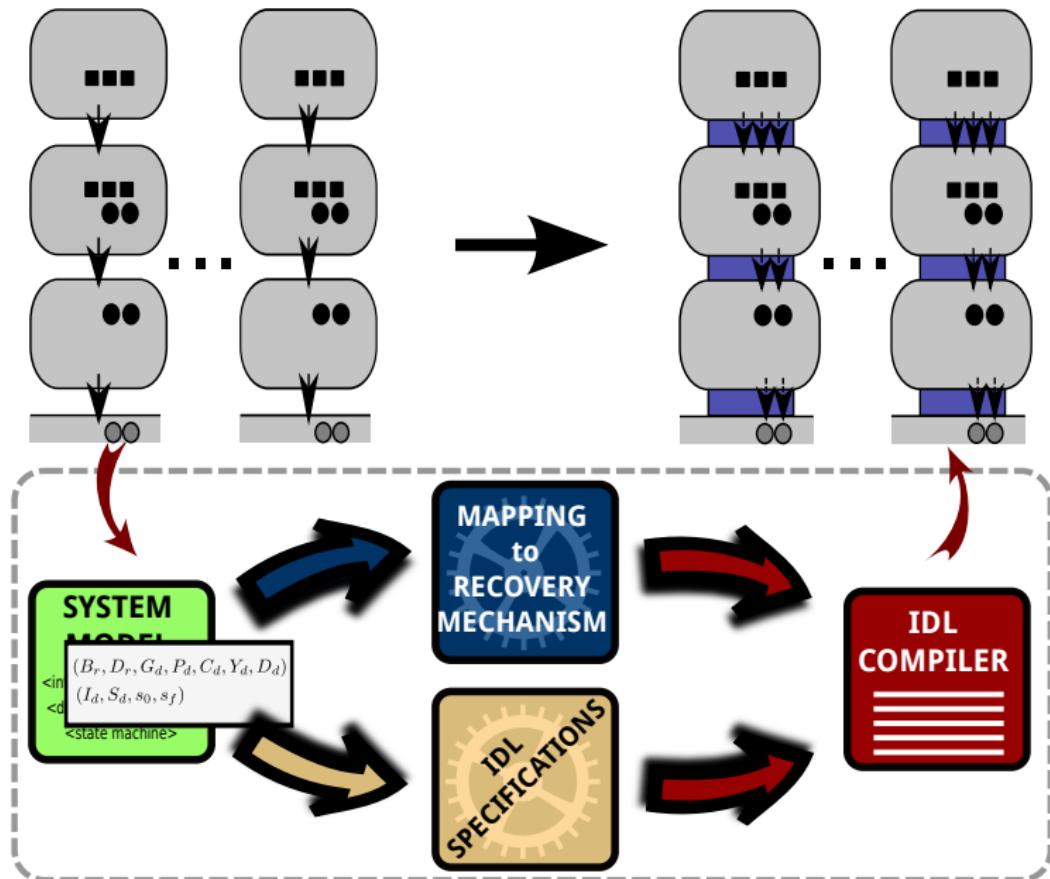
- evaluate predicates
- generate the code from templates



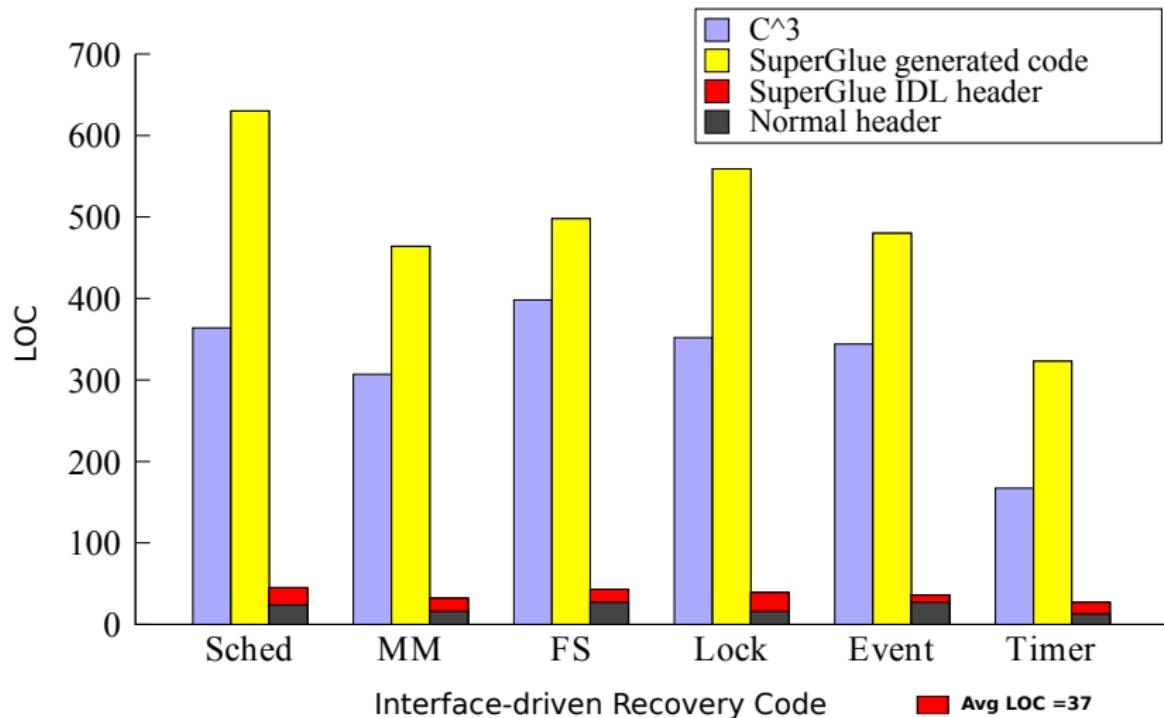
# Generate Recovery Code



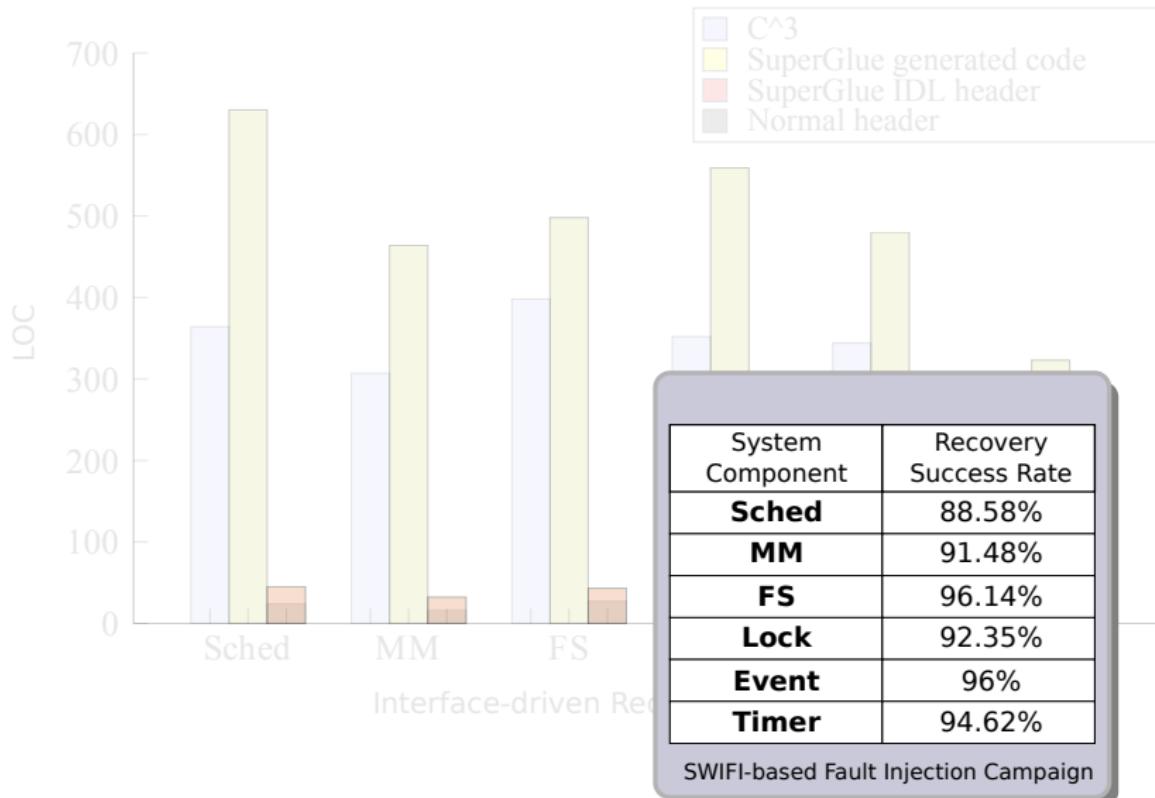
# Generate Recovery Code



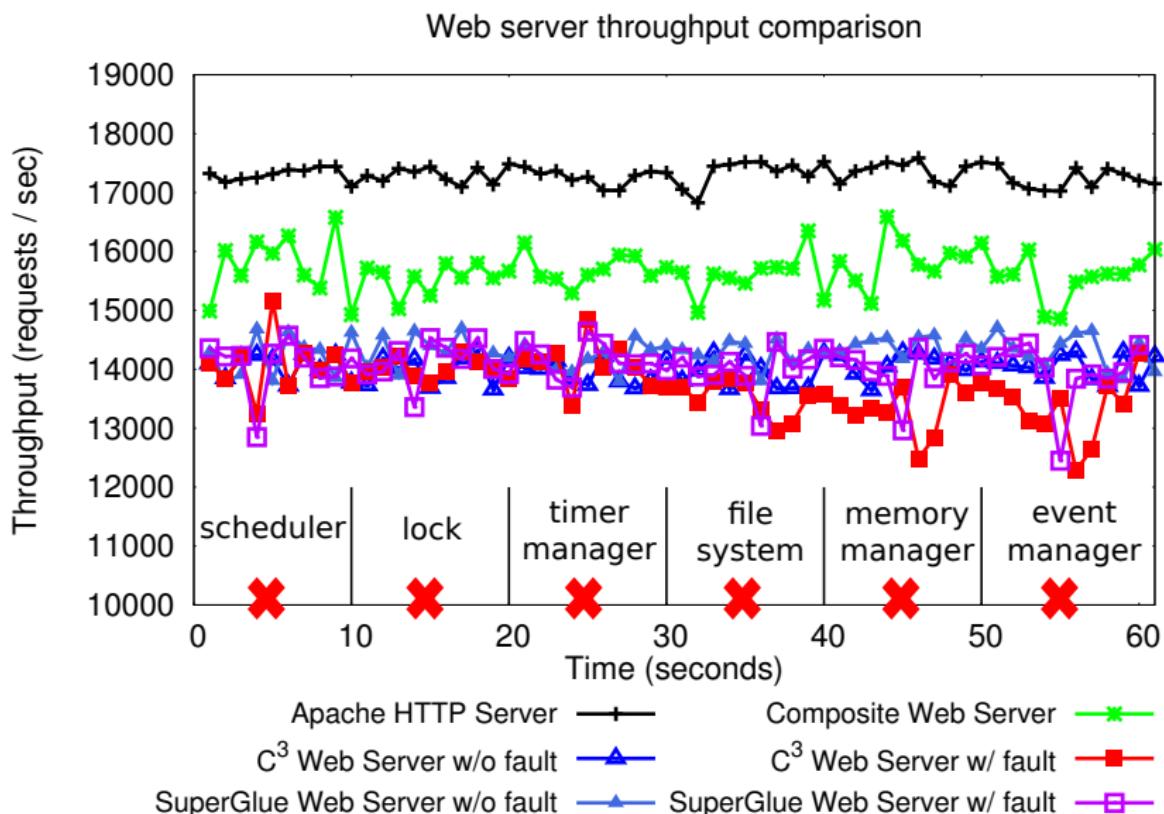
# Code Generation Result



# Fault Injection Result



# Web Server Evaluation with Injected Faults



# Conclusion on SuperGlue

SuperGlue – code generation for system-level fault tolerance

- descriptor-resource **model** and descriptor **state machine**
- **IDL-based** declarative specifications
- **compiler** for synthesizing C<sup>3</sup>-style recovery code

# Conclusion

Thanks

? || /\* \*/

composite.seas.gwu.edu